

3_Interrupts

Motivation

- When a user process is running on the CPU, how does OS stop it?
- When a user process accesses disk, how does OS verify that user has permission?
- When a user process crashes, how does OS regain control?

Topics

- Kernel Mode verses User Mode
- Interrupt Types
- IDT setup
- Interrupt lookup
- HW Interrupt
- System Calls
- Exceptions
- Examples

Kernel Mode verses User Mode

- Kernel Mode
 - Most privileged
 - Access to entire file system, memory space, all hardware
 - OS runs in this mode
- User Mode
 - Least privileged
 - Access to resources (like files and memory) that belong to current user
 - User processes run in this mode

Kernel Mode verses User Mode

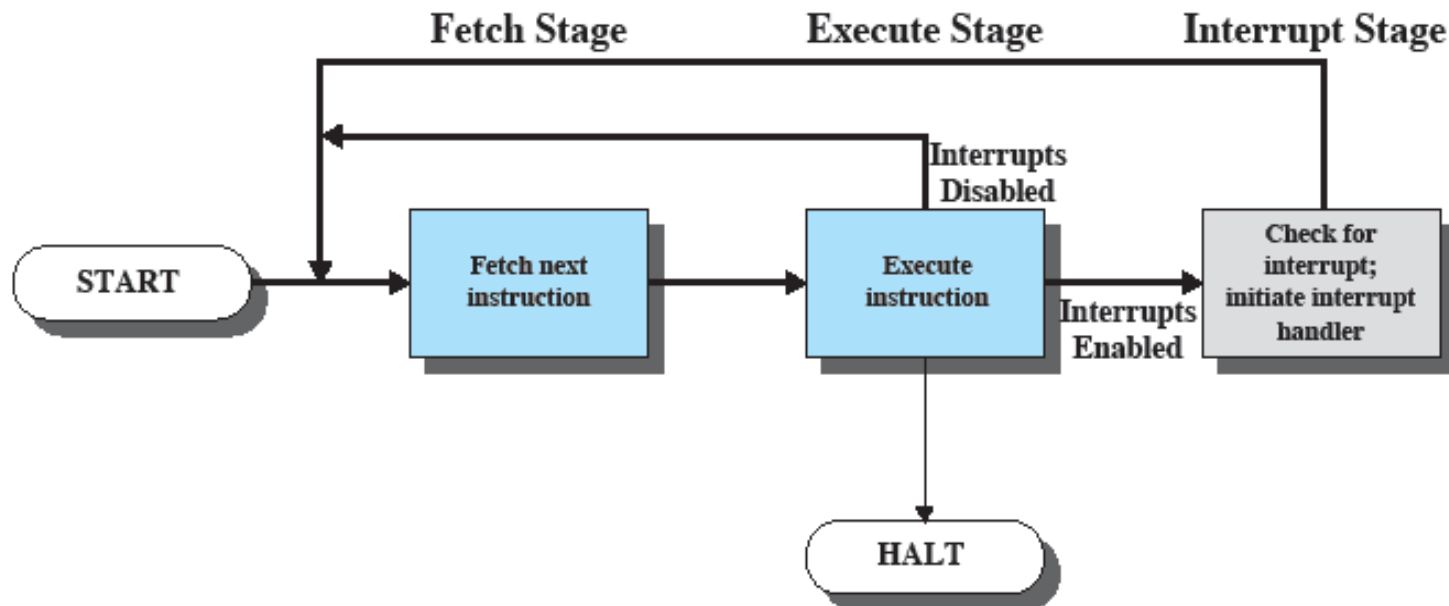
- Can be implemented with 1 bit in a CPU register, register is only accessible in Kernel mode (**means user mode programs cannot modify this bit**)
- What mode is process in?
 - 0=Kernel,
 - 1=User
- Set/Reset when handling interrupts (more coming)

Interrupt Types

- Hardware
 - Raised by hardware devices
 - Can occur at any time (**Asynchronous**)
 - Examples: timer (process switch), I/O signals
- System calls:
 - Software interrupts (**Synchronous**)
 - Raised by user programs to invoke OS functionality
- Exceptions
 - Generated by processor as a result of illegal action (**Synchronous**)
 - **Faults**: recoverable (page fault)
 - **Aborts**: difficult to recover (divide by 0)

Interrupt servicing

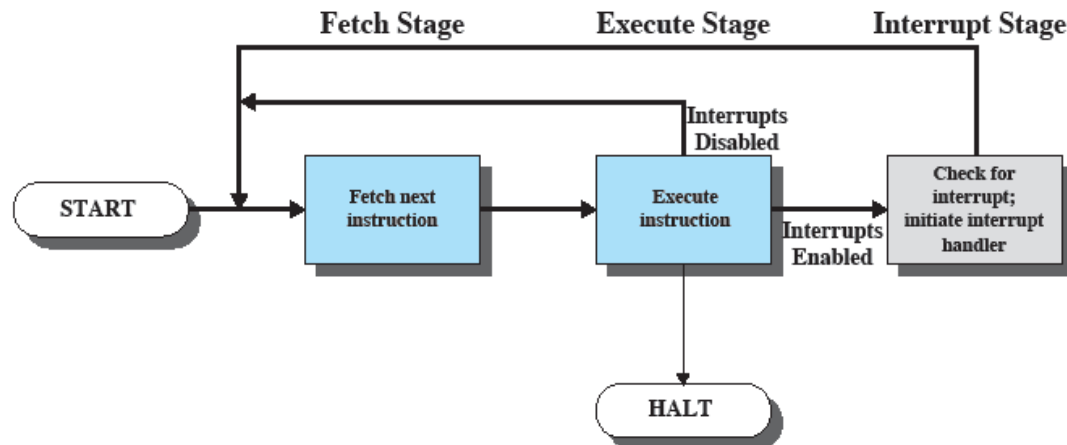
- Remember this cycle is for 1 machine instruction (assembly)
- NOT a high level language (like C++)
- Uses lookup table to find correct interrupt handler (**coming soon!**)
- Works for both synchronous and asynchronous interrupts



Interrupt servicing – pseudo code

```

while (fetch next instruction) {
    run instruction;
    if (interrupt occurred) {
        save process state    // user mode
        find and jump to OS-provided interrupt handler // kernel mode
        restore original process state from kernel stack
    }
    // user mode
}
  
```

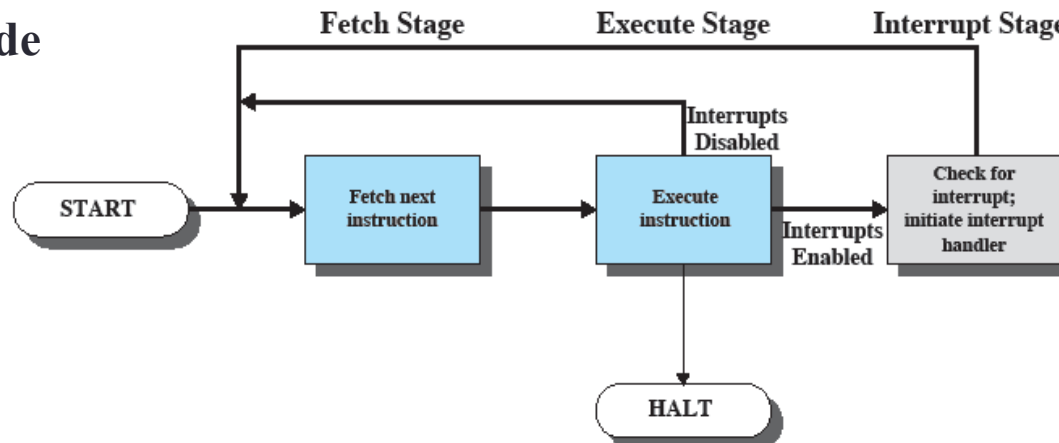


Interrupt servicing – better pseudo code

```

while (fetch next instruction) {
    run instruction;
    if (interrupt occurred) {
        save process state    // user mode
        find and jump to OS-provided interrupt handler // kernel mode
        OS chooses next process to run
        if new process
            load its state from mem
        else
            restore original process state from kernel stack
    }
    // user mode
}

```



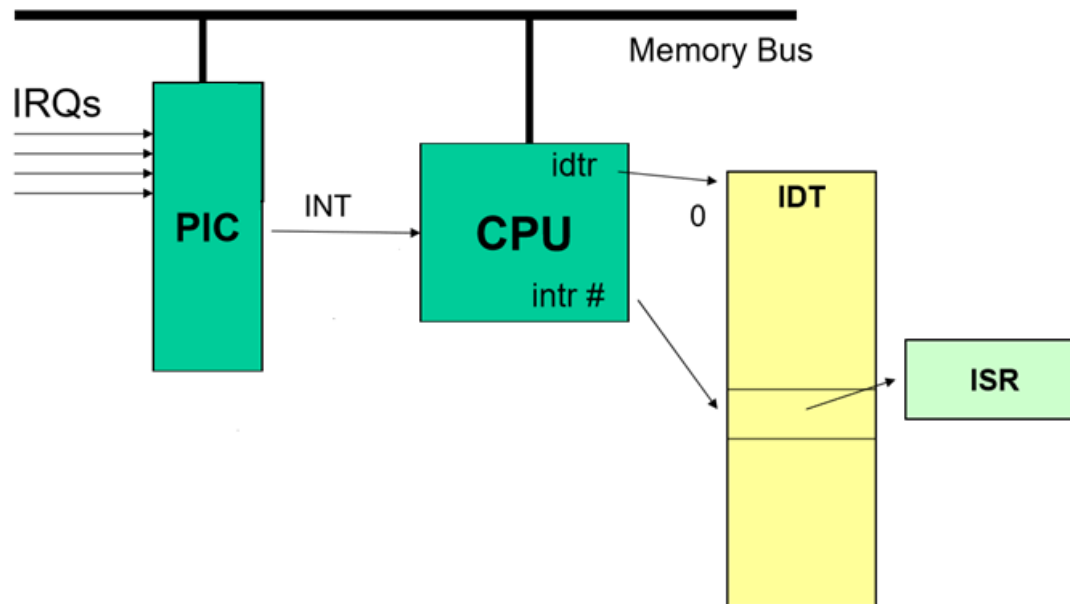
OS kernel
address of display handler address of mouse handler address of disk handler ... address of printer handler
display interrupt handler mouse interrupt handler disk interrupt handler ... printer interrupt handler
window manager
more OS programs
user program 1
user program 2
...
user program n
empty (free) space

- OS sets up IDT at boot, its base pointed to by IDT register (IDTR) in CPU
- Each entry is address of an interrupt handler (known as Interrupt Service Routine (ISR))
- Each ISR handles a particular type of interrupt.

The diagram illustrates the interrupt process. On the left, a vertical teal box labeled **PIC** (Programmable Interrupt Controller) receives multiple horizontal arrows labeled **IRQs** from the left. A horizontal arrow labeled **INT** points from the **PIC** to a teal box labeled **CPU**. The **CPU** box has two labels: **idtr** at the top and **intr #** at the bottom. A horizontal line at the top is labeled **Memory Bus**. An arrow points from the **idtr** label in the **CPU** to a yellow vertical box labeled **IDT** (Interrupt Descriptor Table). The arrow is labeled with the number **0**. The **IDT** box is divided into three horizontal sections. An arrow points from the **intr #** label in the **CPU** to the middle section of the **IDT**. Another arrow points from this middle section to a light green box labeled **ISR** (Interrupt Service Routine).

Hardware Interrupt

Asynchronous interrupt, can happen anytime



- Intr# generated by CPU
- In kernel mode – look up Intr# in IDT
- Run ISR associated with Intr#
- Either return to running process or load and run new one

System Calls

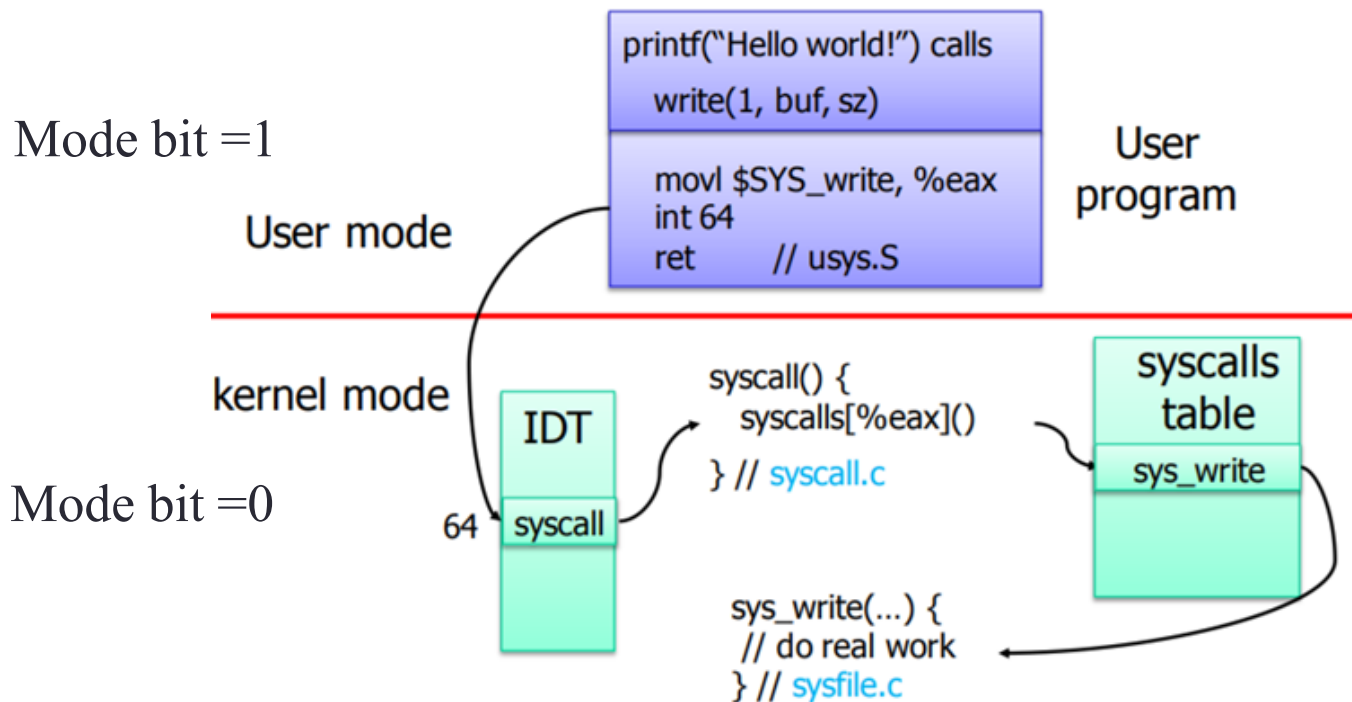
- Applications cannot perform privileged operations themselves
 - Otherwise they could access devices they should not (DMA controller to access other processes memory), write files they don't have permission too, elevate their privileges, etc.
- Instead, they ask OS to do so on their behalf by issuing system calls
- OS verifies permissions, then parameters then fulfills request.

System Call Types

- **Process Control** - process creation, process termination etc.
 - Ex. `fork()`, `exit()`, `wait()`
- **File Management**- file manipulation such as creating a file, reading a file, writing into a file etc.
 - Ex. `open()`, `read()`, `write()`, `close()` //to file
- **Device Management** - device manipulation such as reading from device buffers, writing into device buffers etc.
 - Ex. `read()`, `write()` //to console
- **Information Maintenance** - handle information and its transfer between the operating system and the user program.
 - Ex. `getpid()`, `alarm()`, `sleep()`
- **Communication** - interprocess communication. They also deal with creating and deleting a communication connection.
 - Ex `pipe()`

System Call

synchronous interrupt, you know its coming

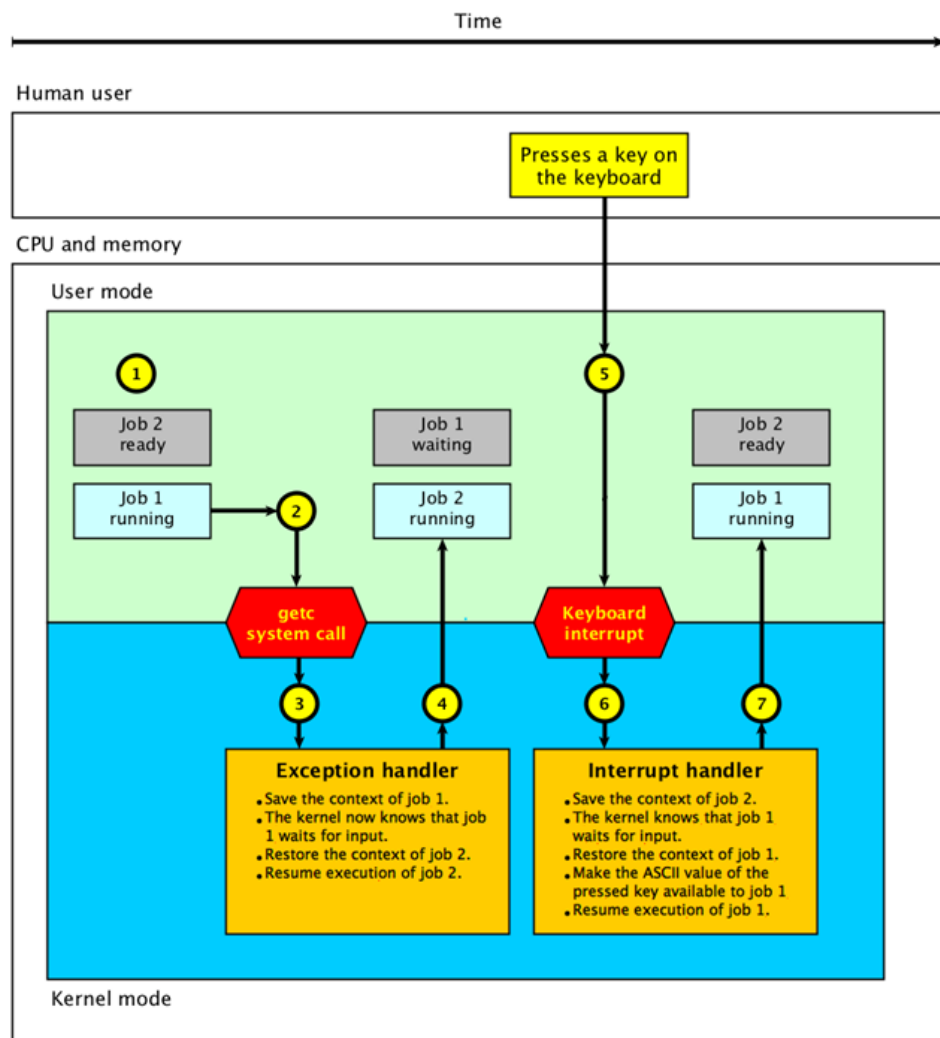


- int 64 passes control to interrupt 64
- In kernel mode – look up interrupt 64 in IDT
- Run `sys_write(...)`
- Either return to running process or load and run new one

Exceptions

- Divide by 0, dereference Null pointer, page fault
- Throws exception interrupt
- Transition to Kernel mode to handle interrupt type as defined in IDT
- Some are recoverable;
 - OS blocks current process, request page from mem, loads another process. When page available OS interrupted again and page is loaded and page fault process can run again
- Some are not
 - Divide by zero, OS logs error, terminates current process, loads a new process

Example - 2 programs, HW and system call interrupts



1. Job 2 is ready to run and job 1 is running.

2. Job 1 uses the `getc` system call to read a character from the keyboard.

3. The system call uses a **system call exception** to enter the kernel.

The kernel saves the context (all register values) of job 1.

The kernel now knows that job 1 waits for input and changes its state from **running** to **waiting**.

The kernel restores the context of job 2.

4. The kernel resumes execution of job 2. Job 1 and changes its state from **ready** to **running**.

5. The human user presses a key on the keyboard.

6. The key-press causes a keyboard **interrupt** which is handled by the kernel.

The kernel saves the context of job 2 and changes its state from **running** to **ready**.

The kernel knows that job 1 waits for the `getc` system call to complete.

The kernel restores the context of job 1 and changes its state from **waiting** to **running**.

The ASCII value of the pressed key is made available to job 1.

7. The kernel resumes execution of job 1.

Summary

- Kernel versus user mode, when and how they are entered
- IDT and ISRs
- Types of interrupts (HW, System call, exception)
- Examples of both hardware and software interrupts

References

- This presentation was developed with the aid of the OSTEP text and the following websites
- http://faculty.salina.k-state.edu/tim/oss/Introduction/sys_calls.html
- <http://www.cs.columbia.edu/~junfeng/13fa-w4118/lectures/l06-trap.pdf>
- <http://www.it.uu.se/education/course/homepage/os/vt18/module-1/system-call-design/>