3_Interrupts

Motivation

- When a user process is running on CPU, how does OS stop it?
- When a user process accesses disk, how does OS verify that user has permission?
- When a user process crashes, how does
 OS regain control?

Topics

- Kernel Mode verses User Mode
- Interrupt Types
- IDT setup
- Interrupt lookup
- HW Interrupt
- System Calls
- Exceptions
- Examples

Kernel Mode verses User Mode

Kernel Mode

- Most privileged
- Access to entire file system, memory space, all hardware
- OS runs in this mode

User Mode

- Least privileged
- Access to resources (like files and memory) that belong to current user
- User processes run in this mode

Kernel Mode verses User Mode

 Implemented with 1 bit in a CPU register, only accessible in Kernel mode (means user mode programs cannot modify this bit)

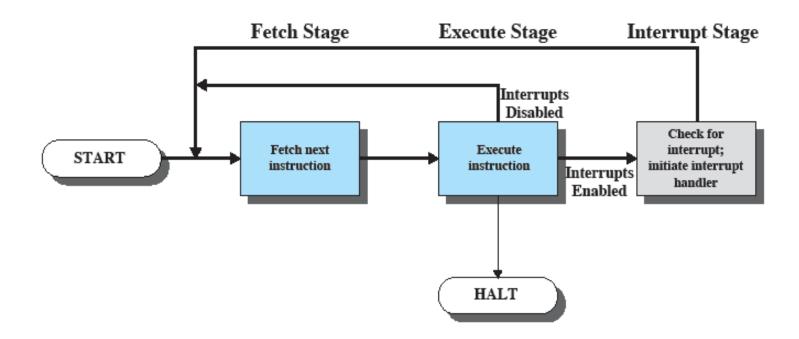
- What mode is process in?
 - 0=Kernel,
 - 1=User
- Set/Reset when handling interrupts (more coming)

Interrupt Types

- Hardware
 - Raised by hardware devices
 - Can occur at any time (Asynchronous)
 - Examples: timer (process switch), I/O signals
- System calls:
 - Software interrupts (Synchronous)
 - Raised by user programs to invoke OS functionality
- Exceptions
 - Generated by processor as a result of illegal action (Synchronous)
 - Faults: recoverable (page fault)
 - Aborts: difficult to recover (divide by 0)

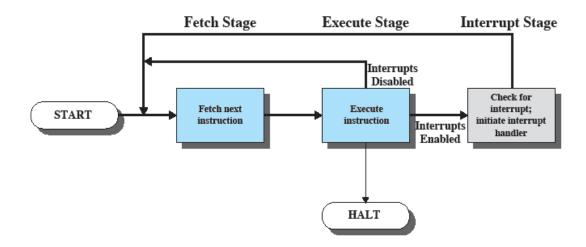
Interrupt servicing

- Remember this cycle is for 1 <u>machine</u> instruction (assembly)
- NOT a high level language (like C++)
- Uses lookup table to find correct interrupt handler (coming soon!)
- Works for both synchronous and asynchronous interrupts



Interrupt servicing – pseudo code

```
while (fetch next instruction) {
    run instruction;
    if (interrupt occurred) {
        save process state // user mode
        find and jump to OS-provided interrupt handler // kernel mode
        restore original process state from kernel stack
    }
    // user mode
}
```



Interrupt servicing – better pseudo code

```
while (fetch next instruction) {
   run instruction;
   if (interrupt occurred) {
       save process state // user mode
       find and jump to OS-provided interrupt handler // kernel mode
       OS chooses next process to run
       if new process
           load its state from mem
       else
          restore original process state from kernel stack
                                   Fetch Stage
                                                    Execute Stage
                                                                     Interrupt Stage
   // user mode
                                                          Interrupts
                                                          Disabled
                                                                         Check for
                                     Fetch next
                                                        Execute
                                                                         interrupt;
                    START
                                     instruction
                                                       instruction
                                                                       initiate interrupt
                                                                Interrupts
                                                                          handler
                                                                 Enabled
                                                        HALT
```

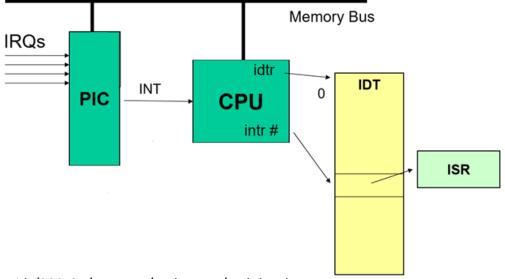
Interrupt Descriptor Table (IDT) setup

IDT Memory Layout



- OS sets up IDT at boot, its base pointed to by IDT register (IDTR) in CPU
- Each entry is address of an interrupt handler (known as Interrupt Service Routine ISR)
- Each ISR handles a particular type of interrupt.

Hardware Interrupt Sequence

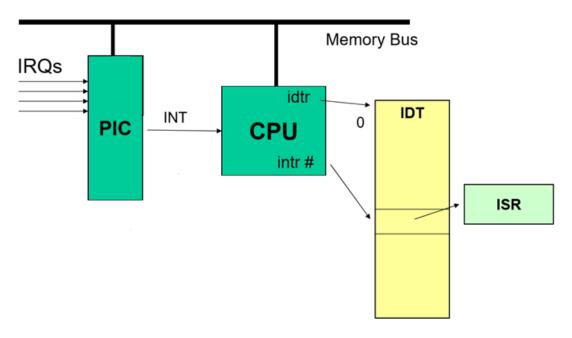


Diagrams from http://people.cs.ksu.edu/~schmidt/300s05/Lectures/ArchNotes/arch.html

interrupt vector

Hardware Interrupt

Asynchronous interrupt, can happen anytime



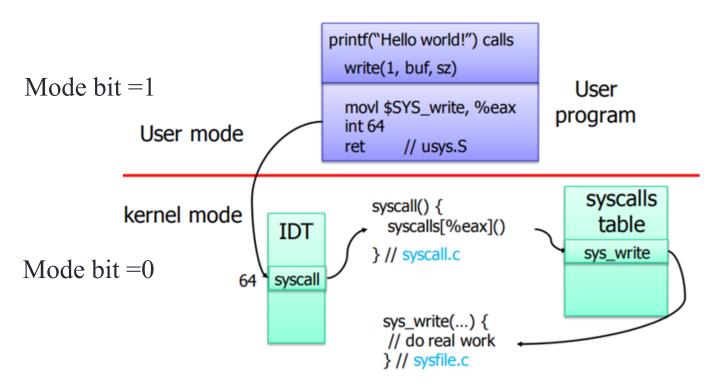
- Intr# generated by CPU
- In kernel mode look up Intr# in IDT
- Run ISR associated with Intr#
- Either return to running process or load and run new one

System Calls

- Applications cannot perform privileged operations themselves
 - Otherwise they could access devices they should not (DMA controller to access other processes memory), write files they don't have permission too, elevate their privileges, etc.
- Instead, they ask OS to do so on their behalf by issuing system calls
- OS verifies permissions, then parameters then fulfills request.

System Call

synchronous interrupt, you know its coming



- int 64 passes control to interrupt 64
- In kernel mode look up interrupt 64 in IDT
- Run sys_write(...)
- Either return to running process or load and run new one

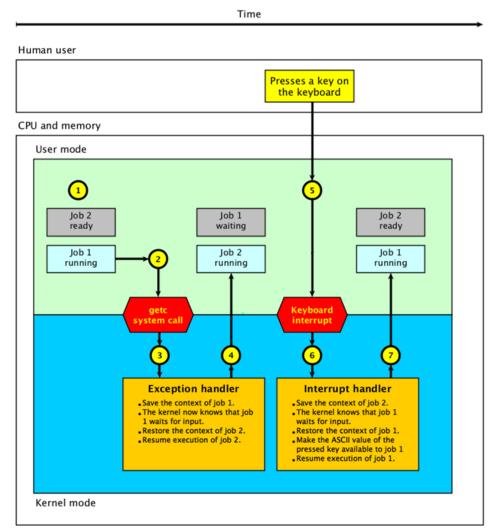
System Call Types

- Process Control process creation, process termination etc.
 - Ex. fork(), exit(), wait()
- File Management- file manipulation such as creating a file, reading a file, writing into a file etc.
 - Ex. open(), read(), write(), close() //to file
- Device Management device manipulation such as reading from device buffers, writing into device buffers etc.
 - Ex. read(), write() //to console
- Information Maintenance handle information and its transfer between the operating system and the user program.
 - Ex. getpid(), alarm(), sleep()
- Communication interprocess communication. They also deal with creating and deleting a communication connection.
 - Ex pipe()

Exceptions

- Divide by 0, dereference Null pointer, page fault
- Throws exception interrupt
- Transition to Kernel mode to handle interrupt type as defined in IDT
- Some are recoverable;
 - OS blocks current process, request page from mem, loadas another process. When page available OS interrupted again and page is loaded and page fault process can run again
- Some are not
 - Divide by zero, OS logs error, terminates current process, loads a new process

Example - 2 programs, HW and system call interrupts



- 1. Job 2 is ready to run and job 1 is running.
- 2. Job 1 uses the getc system call to read a character from the keyboard.
- 3. The system call uses a **system call exception** to enter the kernel.

The kernel saves the context (all register values) of job 1.

The kernel now knows that job 1 waits for input and changes its state from **running** to **waiting**.

The kernel restores the context of job 2.

- 4. The kernel resumes execution of job 2. Job 1 and changes its state from **ready** to **running**.
- 5. The human user presses a key on the keyboard.
- 6. The key-press causes a keyboard **interrupt** which is handled by the kernel.

The kernel saves the context of job 2 and changes its state from **running** to **ready**.

The kernel knows that job 1 waits for the getc system call to complete.

The kernel restores the context of job 1 and changes its state from **waiting** to **running**.

The ASCII value of the pressed key is made available to job 1.

7. The kernel resumes execution of job 1.

Summary

- Kernel verses user mode, when and how they are entered
- IDT and ISRs
- Types of interrupts (HW, System call, exception)
- Examples of both hardware and software interrupts

References

- This presentation was developed with the aid of the OSTEP text and the following websites
- http://faculty.salina.k-state.edu/tim/ossg/Introduction/sys_calls.html
- http://www.cs.columbia.edu/~junfeng/13fa-w4118/lectures/l06-trap.pdf
- http://www.it.uu.se/education/course/homepage/os/vt18/module-1/system-call-design/