



**Department of Physics,  
Computer Science & Engineering**

CPSC 410 – Operating Systems I

# Process Description & Control

Keith Perkins

Adapted from original slides by Dr. Roberto A. Flores  
Also from “CS 537 Introduction to Operating Systems” Arpaci-Dusseau

# Chapter 3 Topics

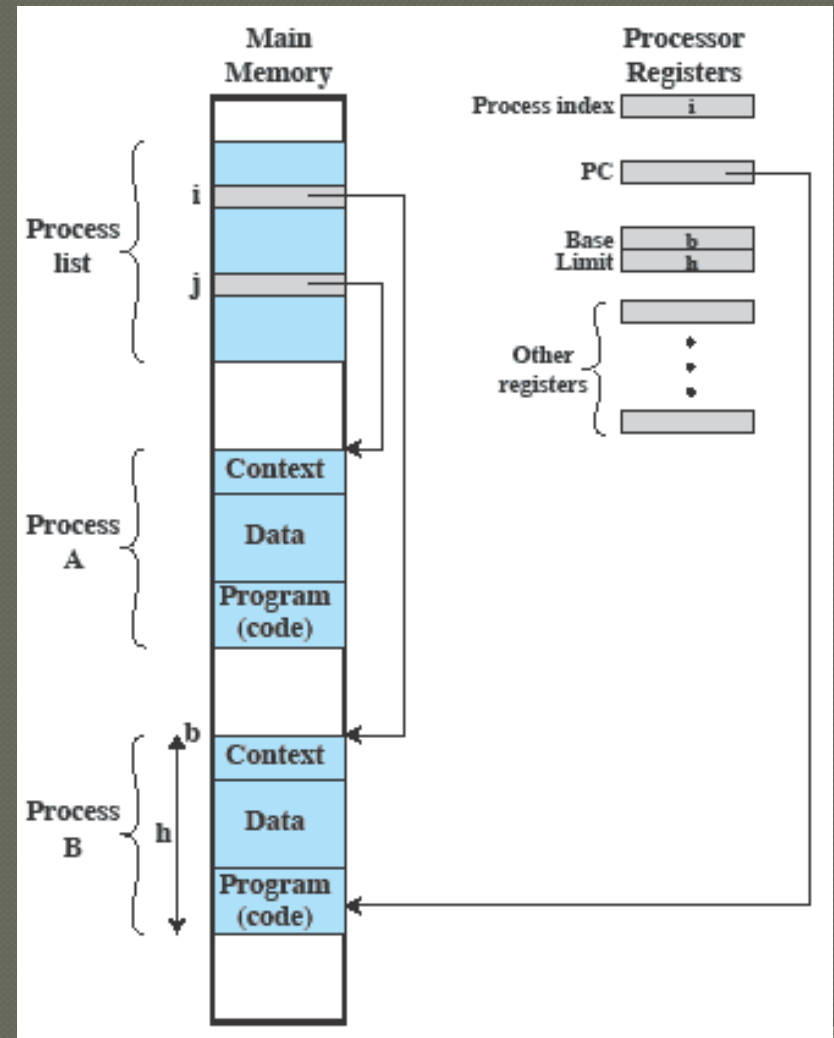
---

## Everything about Processes

- Control blocks
  - States
  - Description
  - Control
- OS Execution
  - Security Issues

# Revisit - Process Management

- Scheduler chooses a process to run (more later)
- Dispatcher runs it
- How? What's in the Process List?
- BTW this list is a simplification



# Processes

## Control Blocks

- data structure created & managed by OS
  - Identifier**: unique ID
  - State**: (e.g., running, blocked)
  - Priority**: relative to other processes
  - Program counter**: address of next instruction
  - Memory pointers**: to code & data
  - I/O status**: I/O in use/pending
  - Accounting**: CPU time used, IDs, ...
- data to hold/restore process state on interrupt/resume
  - key to support multiprocessing

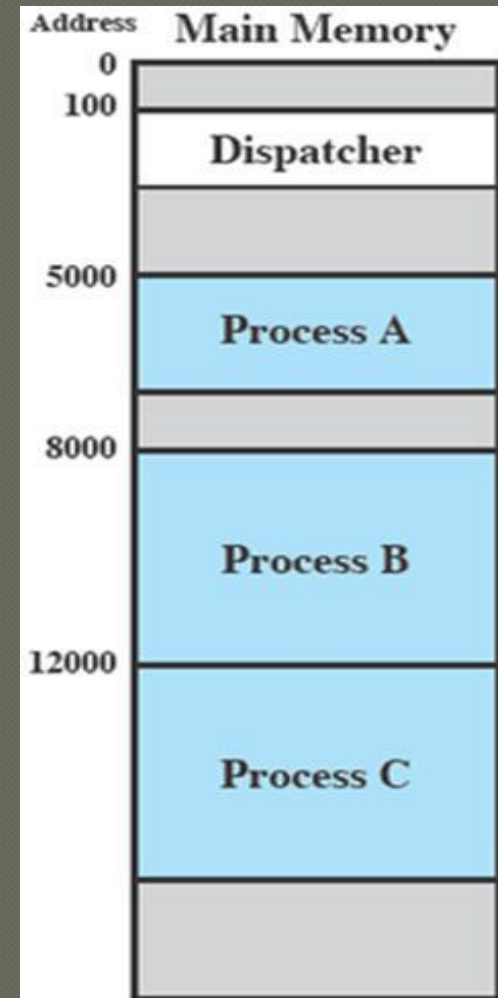
Identifier
State
Priority
Program counter
Memory pointers
Context data
I/O status information
Accounting information
⋮

## Control blocks

States  
Description  
Control

- Dispatcher
  - Program that switches processes in/out of the CPU

# Processes



Control blocks  
States  
Description  
Control

# Processes

## States

- Trace
  - Instructions executed by a process
  - In multiprogramming:
    - interleaving of instructions as processes alternate using the CPU
- The pale blue lower right is dispatcher code
- Process switches because of Interrupts (timer, I/O)

5000	8000	12000
5001	8001	12001
5002	8002	12002
5003	8003	12003
5004		12004
5005		12005
5006		12006
5007		12007
5008		12008
5009		12009
5010		12010
5011		12011

(a) Trace of Process A (b) Trace of Process B (c) Trace of Process C

1	5000	27	12004
2	5001	28	12005
3	5002		
4	5003	29	100
5	5004	30	101
6	5005	31	102
		32	103
		33	104
		34	105
		35	5006
		36	5007
		37	5008
		38	5009
		39	5010
		40	5011
		41	100
		42	101
		43	102
		44	103
		45	104
		46	105
		47	12006
		48	12007
		49	12008
		50	12009
		51	12010
		52	12011

Timeout

Timeout

I/O Request

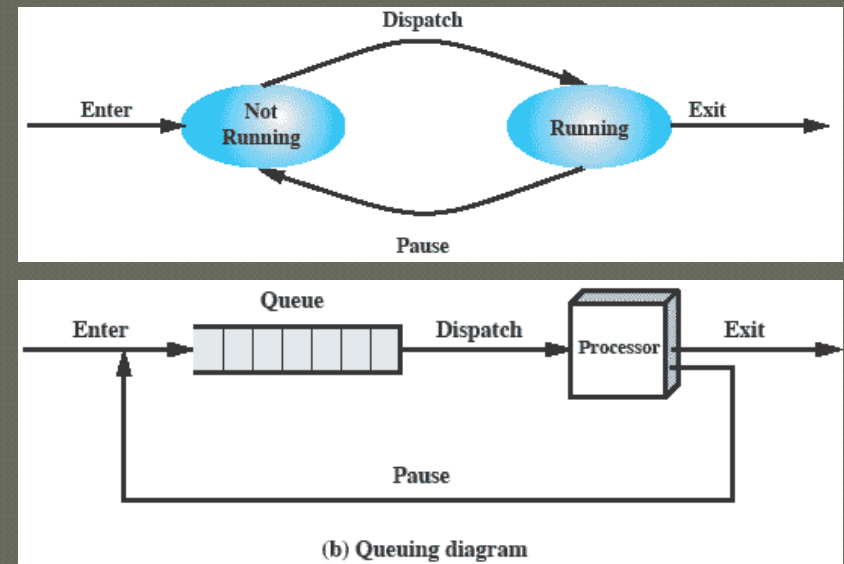
Timeout

Timeout

# Processes

## States (2 states)

- One CPU
- Round-robin (timeout)
- **Running**: CPU time!
- **Not running**: or not



- Where do processes come from?
- When do they stop?


# Processes

- ◎ Where do processes come from? (start)
  - **New batch job**: Next job in the incoming batch stream
  - **Interactive logon**: User in terminal logs in
  - **OS service**: OS-provided service (e.g., print spooler)
  - **Spawned by process**: uses parallelism (parent spawns child)
- ◎ When do they end? (termination)
  - Normal
    - Job finishes, user logs off, OS shutting down, etc.
  - Abnormal
    - **Timeout**: running too long
    - **Resource error**: out of memory, I/O device unresponsive, deadlock
    - **Runtime error**: arithmetic operation, uninitialized variable
    - **Authorization error**: memory out of bounds, resource/instruction privilege




# Processes

## States (5 states)



- **New**: not yet in memory
- **Ready**: awaiting its turn
- **Running**: CPU time!
- **Blocked**: waiting for I/O
- **Exit**: done & gone

# Using Two Queues



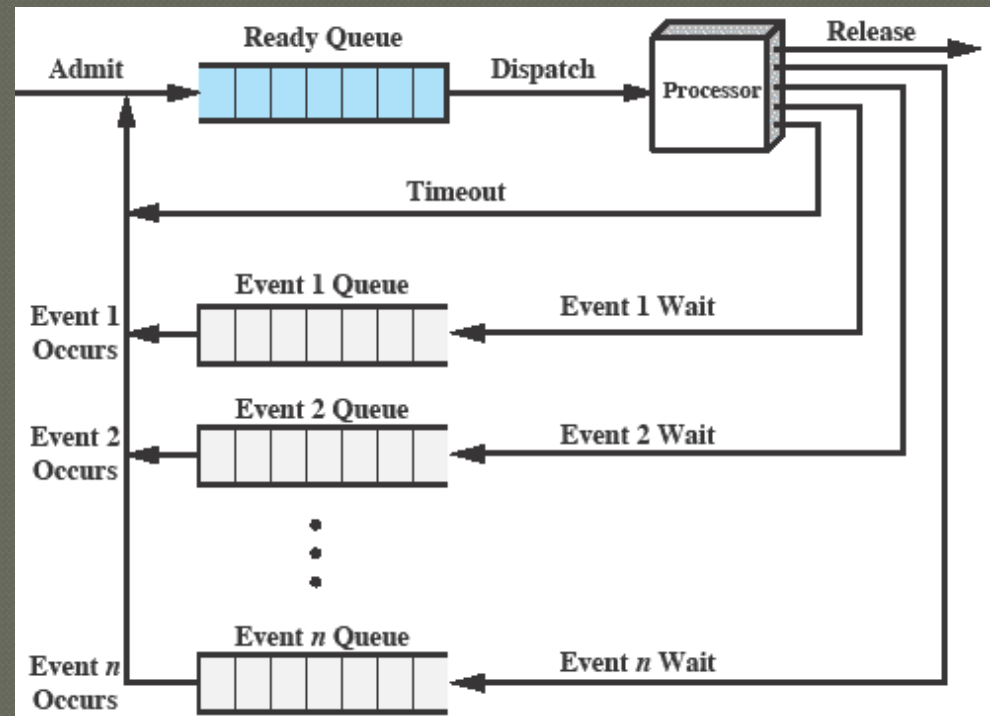

(a) Single blocked queue


Control blocks  
States  
Description  
Control

## States (5 states)

- e.g., Processes A, B & C

- Multiple block queues (1 per I/O device)






# Suspended Processes

- Swapping

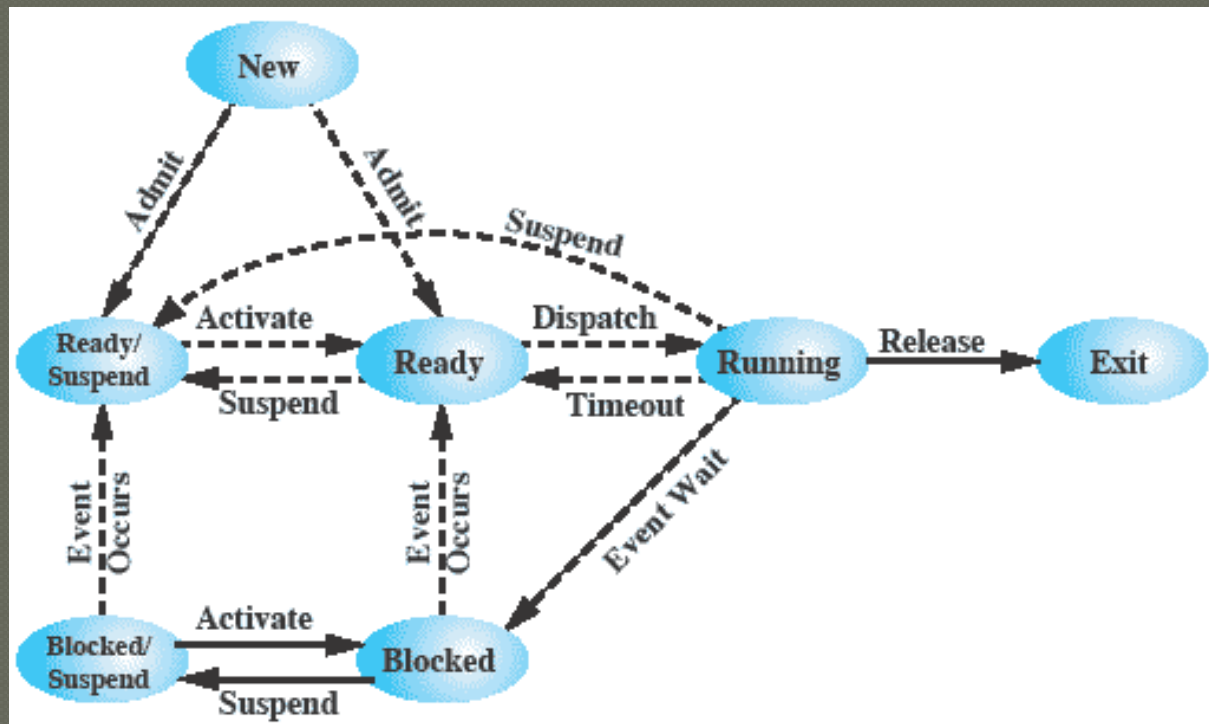
- involves moving part of all of a process from main memory to disk
- when none of the processes in main memory is in the Ready state, the OS swaps one of the blocked processes out on to disk into a suspend queue



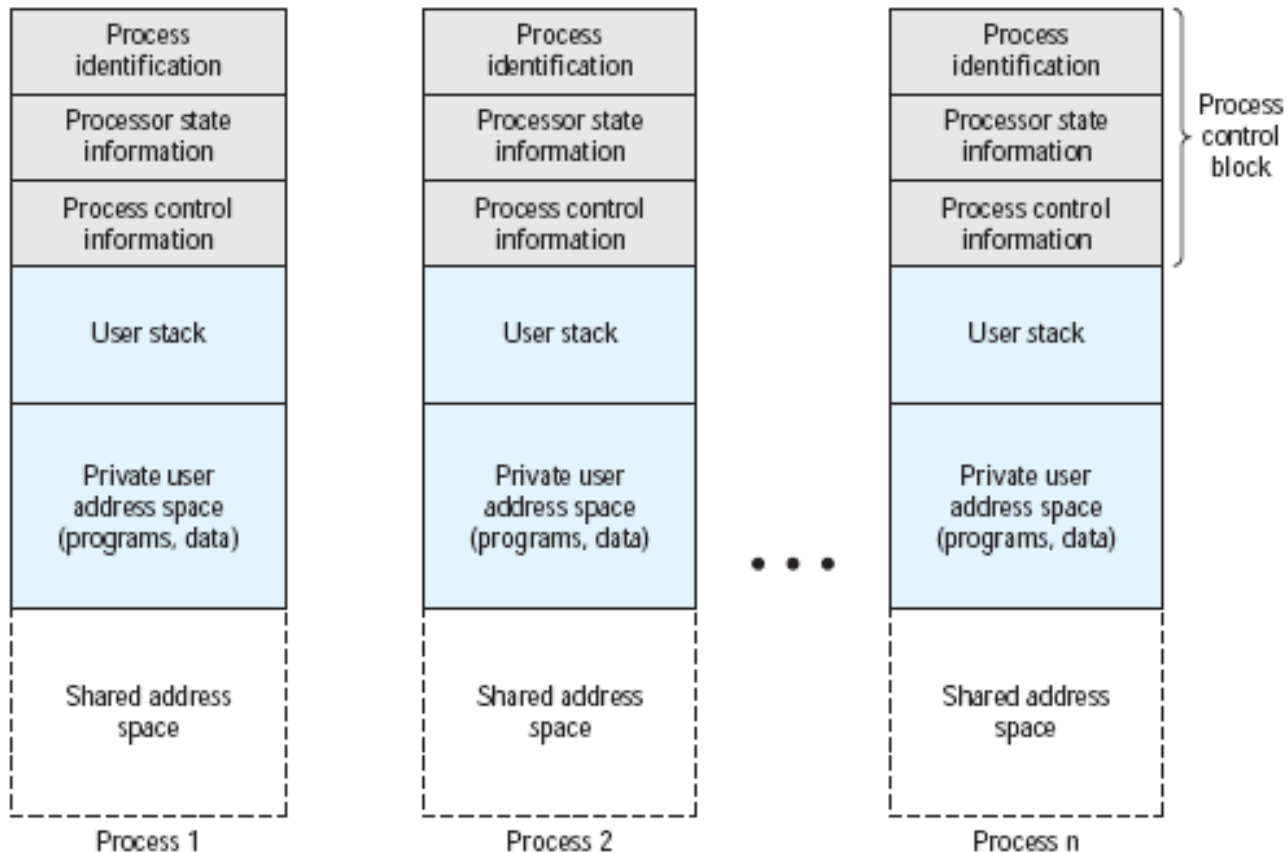
# Processes

## States (7 states)

- What if not all processes fit in memory at once?
- **Suspended**: when a process has been swapped to disk



# Structure of Process Images in Virtual Memory



**Figure 3.13** User Processes in Virtual Memory

# Process List Structures

---

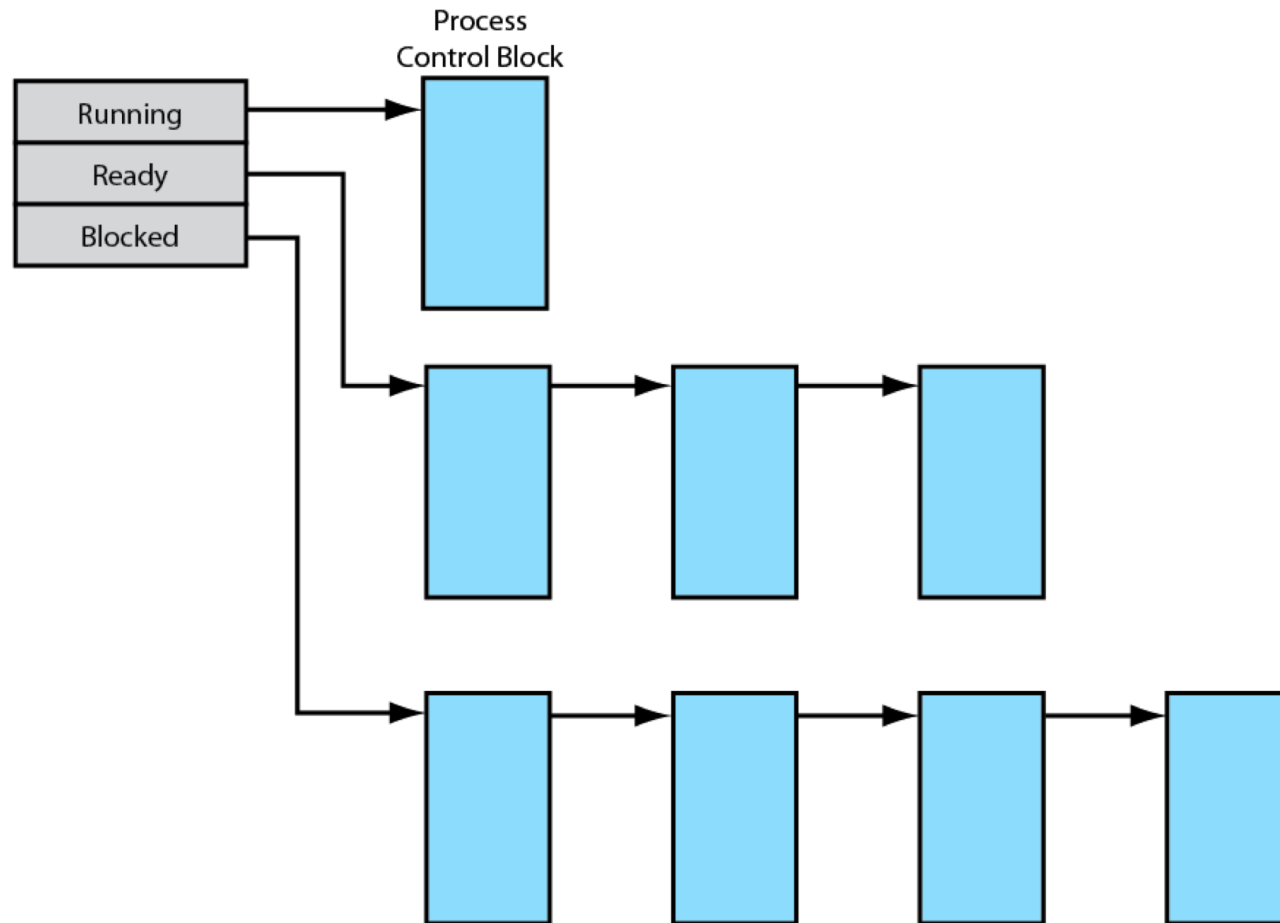


Figure 3.14 Process List Structures

## ● Process tables

- keep data about each process (**process image**)
  - **user data**: modifiable part of program, e.g., variables
  - **user program**: program to execute
  - **stack**: stores method calls & parameters
  - **process control block (PCB)**: data OS uses to control process
    - process **identification**: process/parent/user ID
    - processor **state information**: user/control registers, stack pointers
    - process **control information**: scheduling, inter-process comms, ...
- reference (directly/indirectly) memory, I/O & file tables




## ● Process tables

### Process identification

- Each process has a unique ID
  - IDs are used for reference:
    - in other tables
    - in inter-process communication
    - when a parent spawns a child process
- ➡ process **identification**: process/parent/user ID
  - ➡ processor **state information**: user/control registers, stack pointers
  - ➡ process **control information**: scheduling, inter-process comms, ...
  - reference (directly/indirectly) memory, I/O & file tables

## Process state information

- stack pointers
- user-visible registers
- control & status registers
  - **program status word (PSW)**, e.g., EFLAGS in x86 processors



ID	=	Identification flag
VIP	=	Virtual interrupt pending
VIF	=	Virtual interrupt flag
AC	=	Alignment check
VM	=	Virtual 8086 mode
RF	=	Resume flag
NT	=	Nested task flag
IOPL	=	I/O privilege level
OF	=	Overflow flag

DF = Direction flag  
IF = Interrupt enable flag  
TF = Trap flag  
SF = Sign flag  
ZF = Zero flag  
AF = Auxiliary carry flag  
PF = Parity flag  
CF = Carry flag

- ➡ processor **state information**: user/control registers, stack pointers
- ➡ process **control information**: scheduling, inter-process comms, ...
- reference (directly/indirectly) memory, I/O & file tables

## ● Control

- Modes of execution
  - User mode (-privileged) ... Kernel mode (+privileged)
  - User processes are not allowed to perform;
    - General Memory Access
    - Disk I/O
    - Special x86 instructions like lidt
  - A little background first

# Kernel Mode verses User Mode

---

- Can be implemented with 1 bit
- Kernel Mode
  - Most privileged
  - Access to entire file system, memory space, all hardware
  - OS runs in this mode
- User Mode
  - Least privileged
  - Access to resources (like files and memory) that belong to current user
  - User processes run in this mode
- What mode is process in?
  - Architecture typically supports:
    - Status bit in protected processor register indicating mode
    - Restricts ability to perform certain instructions if not in kernel mode

# Interrupts-types

---

- Types
- Hardware
  - Raised by hardware devices
  - Can occur at any time (**Asynchronous**)
  - Examples: timer (process switch), I/O signals
- Traps:
  - Software interrupts (**Synchronous**)
  - Raised by user programs to invoke OS functionality
- Exceptions
  - Generated by processor as a result of illegal action (**Synchronous**)
  - **Faults**: recoverable (page fault)
  - **Aborts**: difficult to recover (divide by 0)

## ● Control

- Modes of execution
  - User mode (-privileged) ... Kernel mode (+privileged)
  - User processes are not allowed to perform;
    - General Memory Access
    - Disk I/O
    - Special x86 instructions like lidt



So how does a process perform restricted operation (like read/write to disk)?

by trapping into the OS and running a handler in kernel mode  
Address of handler is IDT + interrupt number

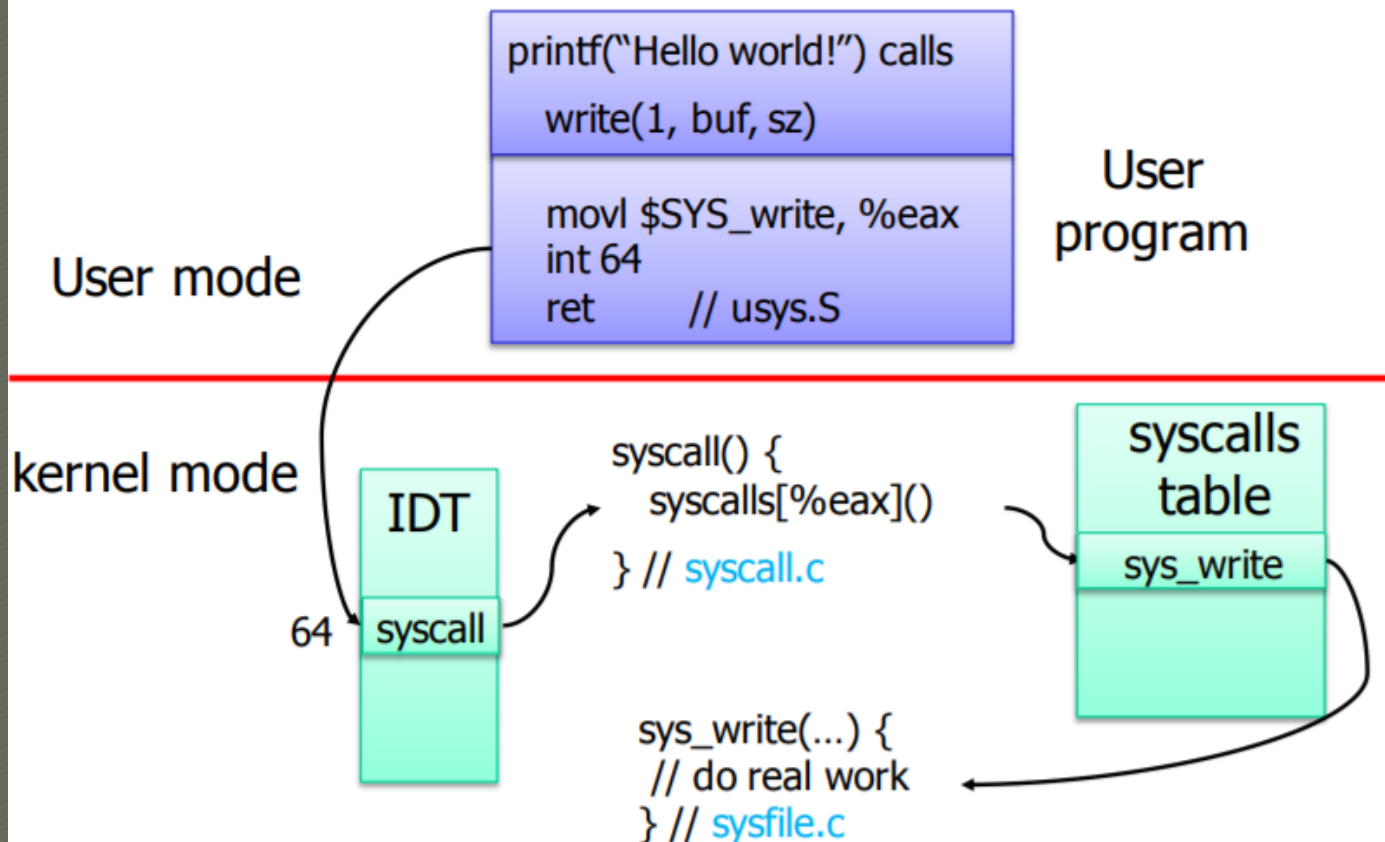
OS sets up an Interrupt Descriptor Table (IDT) at boot (in memory, its base pointed to by IDT register(IDTR) in CPU)  
Each entry is address of an interrupt handler (known as Interrupt Service Routine ISR)  
Each ISR handles a particular type of interrupt.

### System call dispatch

1. Kernel assigns system call type a **system call number**
2. Kernel initializes **system call table**, mapping system call number to function implementing the system call
  - Also called **system call vector**
3. User process sets up system call number and arguments
4. User process runs **int X**
5. Hardware switches to kernel mode and invokes kernel's interrupt handler for **X (interrupt dispatch)**
6. Kernel looks up syscall table using system call number
7. Kernel invokes the corresponding function
8. Kernel returns by running **iret (interrupt return)**




### xv6 system call dispatch



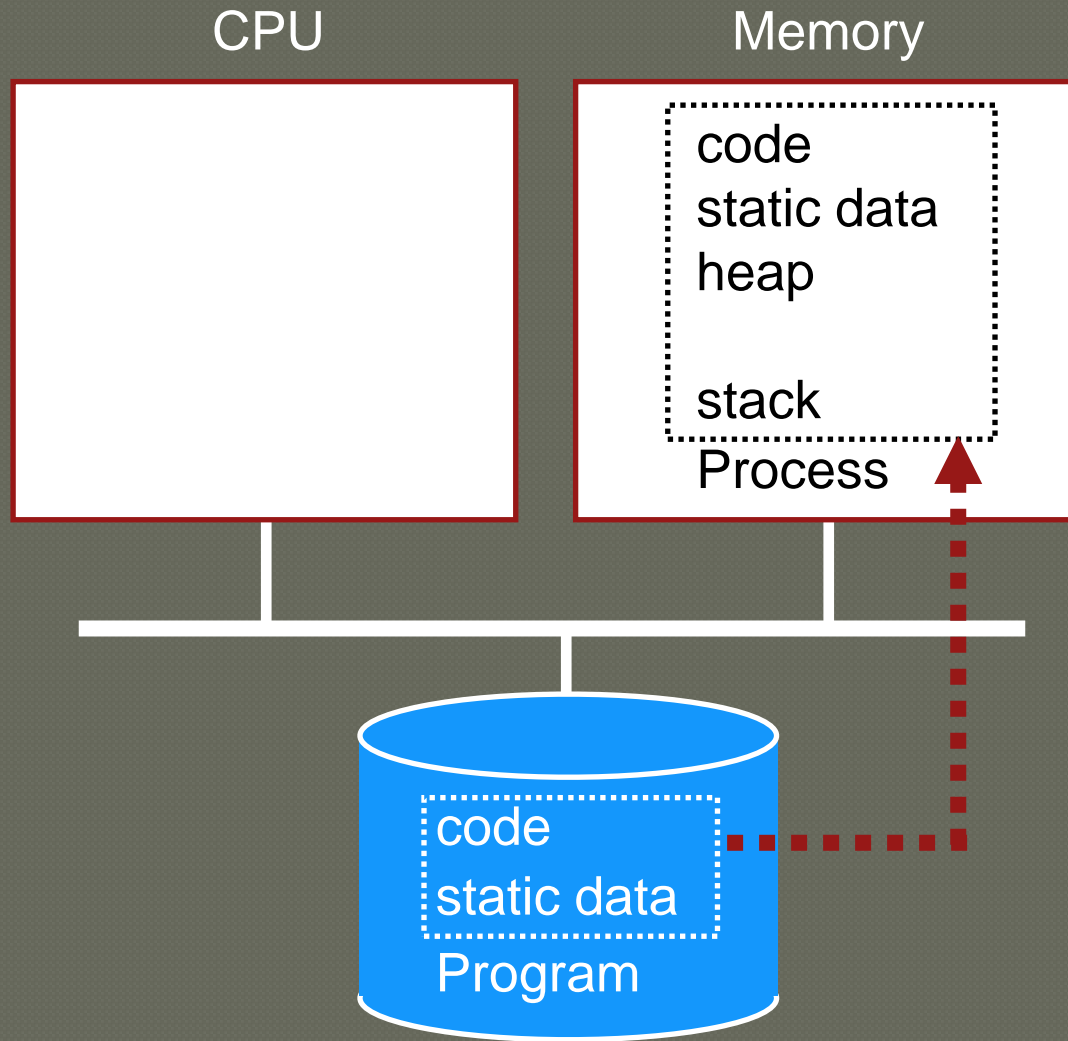
## ● Control

- Process creation
  - What does OS do when a process is created?
    - assigns a new unique ID
    - allocates space for the process in memory
    - initializes its process control block & sets it in place (e.g. in process list)

# Process Creation



# Process Creation



Control blocks  
States  
Description  
Control

# Processes

## Dispatch Mechanism

Process is running- how to switch to other process?

Control blocks  
States  
Description  
Control

# Processes

## Dispatch Mechanism

Process is running- how to switch to other process?

OS runs **dispatch loop**

```
while (1) {  
    run process A for some time-slice  
    stop process A and save its context  
    load context of another process B  
}
```

**Context-switch**

Question 1: How does dispatcher gain control?

Question 2: What execution context must be saved and restored?

Q1: How does Dispatcher get CONTROL?

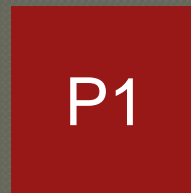
## Option 1: Cooperative Multi-tasking

- Trust process to relinquish CPU to OS through traps
  - Examples: System call, page fault (access page not in main memory), or error (illegal instruction or divide by zero)
  - Provide special `yield()` system call

# Cooperative Approach

---

P1




yield() call




# Cooperative Approach

---



# Cooperative Approach

---



# Cooperative Approach

---

P2

↑ yield() return

# Cooperative Approach

---

P2



yield() call

# Processes

Q1: How does Dispatcher get CONTROL?

- Problem with cooperative approach? **YES**
- Disadvantages: Processes can misbehave
  - By avoiding all traps and performing no I/O, can take over entire machine
  - Only solution: Reboot (windows 95)!
- **Not performed in modern operating systems**

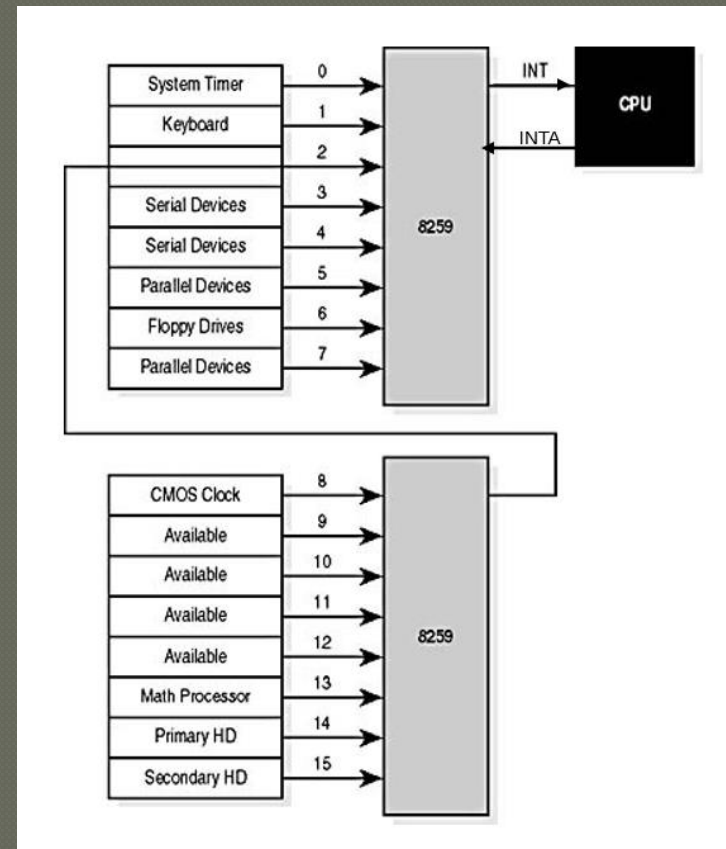
Q1: How does Dispatcher get CONTROL?

## Option 2: True Multi-tasking

- Guarantee OS can obtain control periodically
- Enter OS by enabling periodic alarm clock
  - Hardware generates timer interrupt (CPU or separate chip)
  - Example: Every 10ms
- User must not be able to mask timer interrupt
- Dispatcher counts interrupts between context switches
  - Example: Waiting 20 timer ticks gives 200 ms time slice
  - Common time slices range from 10 ms to 200 ms

# Interrupts-HW- timer example

- 8259 (Programmable interrupt controller or PIC) relays up to 8 interrupt to CPU
- Devices raise interrupts by an 'interrupt request' (IRQ)
- CPU acknowledges and queries the 8259 to determine which device interrupted (int#)
- Priorities can be assigned to each IRQ line
- 8259s can be cascaded to support more interrupts



# Processes

What context to save?

Dispatcher must track context of process when not running

- Save context in **process control block (PCB)**

What information is stored in PCB?

- PID
- Process state (i.e., running, ready, or blocked)
- Execution state (all registers, PC, stack ptr)
- Scheduling priority
- Accounting information (parent and child processes)
- Credentials (which resources can be accessed, owner)
- Pointers to other allocated resources (e.g., open files)

Requires special hardware support

- Hardware saves process PC and PSR on interrupts



Operating System

Hardware

Program

Process A

\*\*\*

Operating System

Hardware

Program

Process A

...

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler

## Operating System

## Hardware

## Program

Process A

\*\*\*

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler

Handle the trap  
Call **switch()** routine  
save regs(A) to proc-struct(A)  
restore regs(B) from proc-struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

## Operating System

## Hardware

## Program

Process A

\*\*\*

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler

restore regs(B) from k-stack(B)  
move to user mode  
jump to B's IP

Handle the trap  
Call **switch()** routine  
save regs(A) to proc-struct(A)  
restore regs(B) from proc-struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

## Operating System

## Hardware

## Program

Process A

...

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler


Handle the trap  
Call **switch()** routine  
save regs(A) to proc-struct(A)  
restore regs(B) from proc-struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

restore regs(B) from k-stack(B)  
move to user mode  
jump to B's IP

Process B

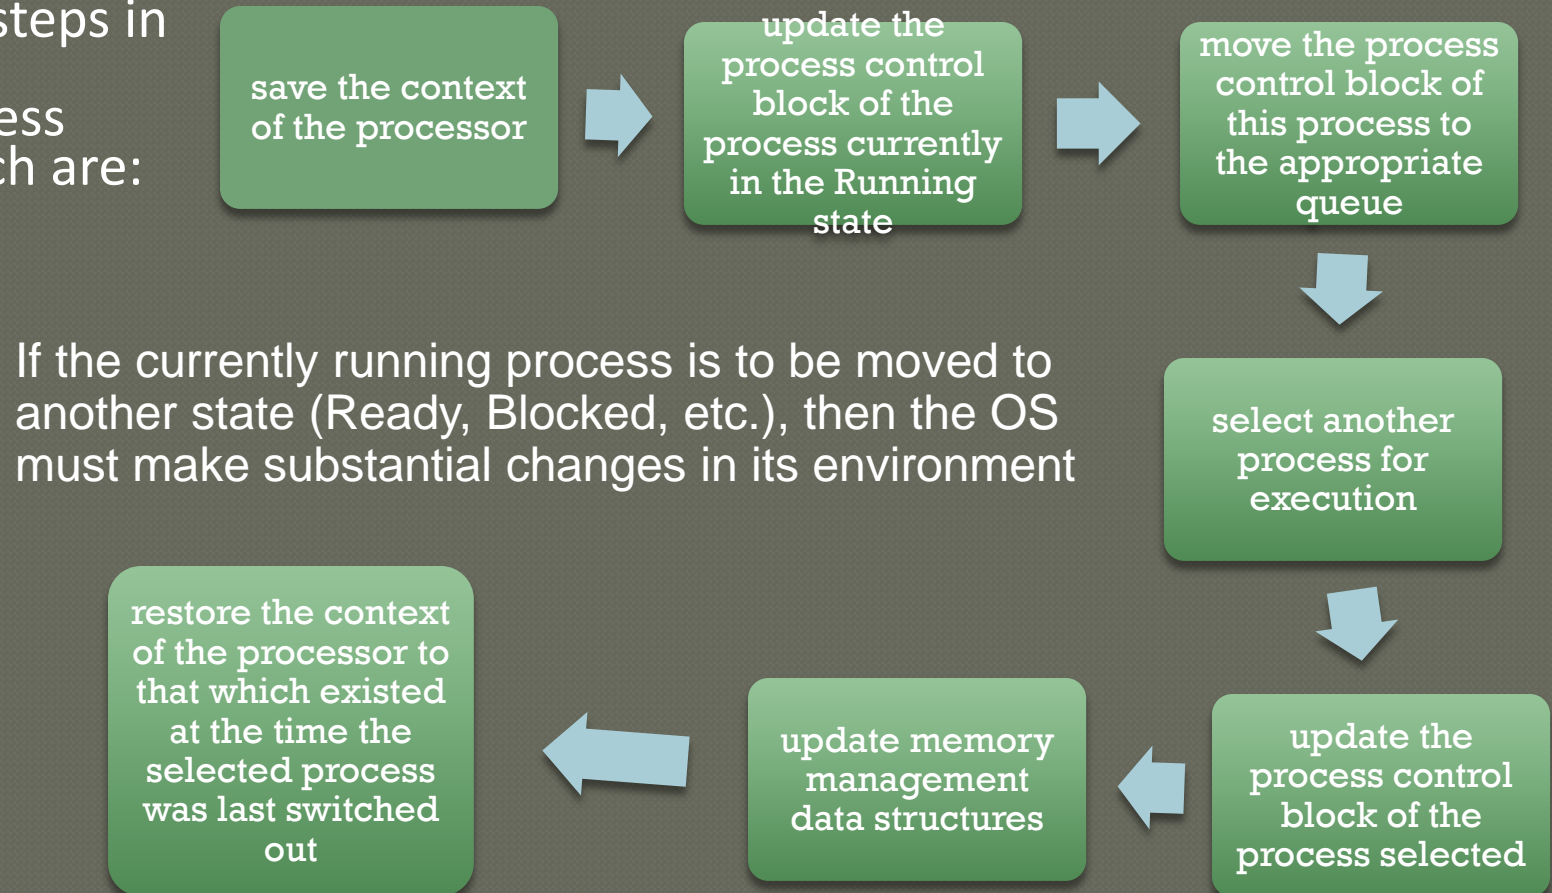
...

# Interrupts



# Change of Process State

- The steps in a full process switch are:



# Mode Switching

If no interrupts are pending the processor:



proceeds to the fetch stage and fetches the next instruction of the current program in the current process

If an interrupt is pending the processor:



sets the program counter to the starting address of an interrupt handler program



switches from user mode to kernel mode so that the interrupt processing code may include privileged instructions



# Chapter 3 Topics

---

- Everything about Processes

- Elements
- Control blocks
- States
- Description
- Control

- OS Execution

- Security Issues

# OS Execution

## ● OS is software, right?

- How is it **different** from just **another process**?
- How is it controlled?

### a) Non-process Kernel


- Processes are processes.  
The kernel is the kernel.

### b) Execution within user processes

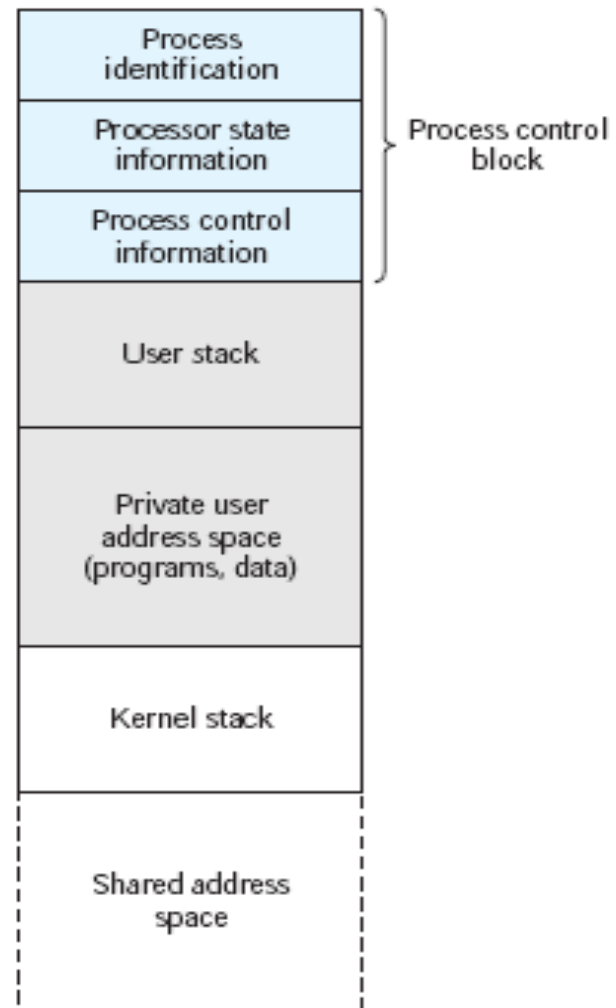
- OS is a bare process switching mechanism
- OS routines are linked to user programs (OS data is shared)

### c) Process-based OS

- OS routines run as independent processes
- Modular approach for parallelism (e.g., OS in one CPU, user processes in another)



# Execution *Within* User Processes



**Figure 3.16** Process Image: Operating System Executes within User Space

# Chapter 3 Topics

---

## ◉ Everything about Processes

- Elements
- Control blocks
- States
- Description
- Control

## ◉ OS Execution

## ◉ Security Issues

# Security

---

## ● Protecting computer resources

- OS should **prevent** (or at least **detect**) users/malware attempts to gain unauthorized access
- **Privileges**
  - Users have privilege levels (highest: administrator/root)
  - Processes have (at most) the same privilege as their user

## ● Threats

- A **potential violation of security**, given a circumstance/capability/action/event breaching security and causing harm.

## ● Countermeasures

- An action/technique that **eliminates/prevents/minimizes/reports** a threat.

## ⦿ Threats

- Goal: gain access to / increase privileges in system
- Intruders (hacker | cracker)
  - **Misfeasor**: user seeking more than allowed | misusing resources
  - **Masquerader**: non-user posing as legitimate user
  - **Clandestine user**: (non-) user seeking root privilege
- Malicious software (malware)
  - Sophisticated (harmless -> crippling)
  - **Parasitic** (needs host program)
    - virus: self-replicating code embedded into another program
    - logic bomb: routine activated under certain conditions
    - backdoor: non-regular access to system (left by designers)
  - **Independent**: worm (virus-minus-host)

## ◉ Countermeasures

- Intrusion detection
  - Service **monitoring** system events, warning about attempts to access resources in an unauthorized manner.
  - 3 logical components
    - sensing >> analyzing >> reporting (UI)
- Authentication
  - Process of verifying an identity claimed by a system entity.
    - **Identification**: representative token
    - **Verification**: examining token
- Firewalls
  - Computer controlling network traffic (based on policies)

# Chapter 3 Topics

---

- Everything about Processes

- Elements
- Control blocks
- States
- Description
- Control

- OS Execution

- Security Issues



Done!