

**Department of Physics,  
Computer Science & Engineering**

CPSC 410 – Operating Systems I

# Virtualization: CPU Scheduling

Keith Perkins

Adapted from “CS 537 Introduction to Operating Systems”  
Arpaci-Dusseau

# CPU Virtualization:

- 
- ◎ Questions answered in this lecture:
    - What are different scheduling policies, such as: FCFS, SJF, STCF, RR and MLFQ?
    - What type of workload performs well with each scheduler?

# CPU Virtualization:

---

## Dispatcher (Previous lecture)

- Low-level mechanism
- Performs context-switch
  - Switch from user mode to kernel mode
  - Save execution state (registers) of old process in PCB
  - Insert PCB in ready queue
  - Load state of next process from PCB to registers
  - Switch from kernel to user mode
  - Jump to instruction in new user process

## ● Scheduler (Today)

- Policy to determine which process gets CPU when

# CPU Virtualization: Two Components

---


## Dispatcher (Previous lecture)

- Low-level mechanism
- Performs context-switch
  - Switch from user mode to kernel mode
  - Save execution state (registers) of old process in PCB
  - Insert PCB in ready queue
  - Load state of next process from PCB to registers
  - Switch from kernel to user mode
  - Jump to instruction in new user process

## ● Scheduler (Today)

- Policy to determine which process gets CPU when

# Review: State Transitions



How to transition? ("mechanism")  
When to transition? ("policy")

# Vocabulary

---

**Workload:** set of **job** descriptions (arrival time, run\_time)

- Job: View as current CPU burst of a process
- Process alternates between CPU and I/O  
process moves between ready and blocked queues

**Scheduler:** logic that decides which ready job to run

**Metric:** measurement of scheduling quality



# Scheduling Performance Metrics

---

## Minimize turnaround time

- Do not want to wait long for job to complete
- $\text{Completion\_time} - \text{arrival\_time}$

## Minimize response time

- Schedule interactive jobs promptly so users see output quickly
- $\text{Initial\_schedule\_time} - \text{arrival\_time}$

## Minimize waiting time

- Do not want to spend much time in Ready queue

## Maximize throughput

- Want many jobs to complete per unit of time

## Maximize resource utilization

- Keep expensive devices busy

## Minimize overhead

- Reduce number of context switches

## Maximize fairness

- All jobs get same amount of CPU over some time interval

# Workload Assumptions

---

1. Each job runs for the same amount of time
2. All jobs arrive at the same time
3. All jobs only use the CPU (no I/O)
4. Run-time of each job is known



# Scheduling Basics

---

## Workloads:

arrival\_time  
run\_time

## Schedulers:

FIFO  
SJF  
STCF  
RR

## Metrics:

turnaround\_time  
response\_time

# Example: workload, scheduler, metric

| JOB | arrival_time<br>(s) | run_time<br>(s) |
|-----|---------------------|-----------------|
| A   | ~0                  | 10              |
| B   | ~0                  | 10              |
| C   | ~0                  | 10              |

**FIFO:** First In, First Out

- also called FCFS (first come first served)
- run jobs in *arrival\_time* order


What is our turnaround?  $\text{completion\_time} - \text{arrival\_time}$

# FIFO: Event Trace

| JOB | arrival_time<br>(s) | run_time<br>(s) | Time | Event      |
|-----|---------------------|-----------------|------|------------|
| A   | ~0                  | 10              | 0    | A arrives  |
| B   | ~0                  | 10              | 0    | B arrives  |
| C   | ~0                  | 10              | 0    | C arrives  |
|     |                     |                 | 0    | run A      |
|     |                     |                 | 10   | complete A |
|     |                     |                 | 10   | run B      |
|     |                     |                 | 20   | complete B |
|     |                     |                 | 20   | run C      |
|     |                     |                 | 30   | complete C |

# FIFO (Identical JOBS)


| JOB | arrival_time<br>(s) | run_time<br>(s) |
|-----|---------------------|-----------------|
| A   | ~0                  | 10              |
| B   | ~0                  | 10              |
| C   | ~0                  | 10              |



Gantt chart:

Illustrates how jobs are scheduled over time on a CPU


# FIFO (IDENTICAL JOBS)



What is the average turnaround time?

Def:  $turnaround\_time = completion\_time - arrival\_time$

# FIFO (IDENTICAL Jobs)



What is the average turnaround time?

Def:  $turnaround\_time = completion\_time - arrival\_time$

$$(10 + 20 + 30) / 3 = 20s$$



# Scheduling Basics

---

## Workloads:

arrival\_time

run\_time

## Schedulers:

FIFO

SJF

STCF

RR

## Metrics:

turnaround\_time

response\_time

# Workload Assumptions

---

- ~~1. Each job runs for the same amount of time~~
2. All jobs arrive at the same time
3. All jobs only use the CPU (no I/O)
4. The run-time of each job is known

# Any Problematic Workloads for FIFO?

---

**Workload:** ?

**Scheduler:** FIFO

**Metric:** turnaround is high

# Example: Big First Job

| JOB | arrival_time<br>(s) | run_time<br>(s) |
|-----|---------------------|-----------------|
| A   | ~0                  | 60              |
| B   | ~0                  | 10              |
| C   | ~0                  | 10              |

Draw Gantt chart for this workload and policy...  
What is the average turnaround time?

# Example: Big First Job

| JOB | arrival_time<br>(s) | run_time<br>(s) |
|-----|---------------------|-----------------|
|-----|---------------------|-----------------|

A ~0 60


B ~0 10

C ~0 10

A: 10  
A: 60s

B: 70s

C: 80s



Average turnaround time: **70s**



# Convoy Effect





# Passing the Tractor

---

## **Problem with Previous Scheduler:**

FIFO: Turnaround time can suffer when short jobs must wait for long jobs

## **New scheduler:**

SJF (Shortest Job First)


Choose job with smallest *run\_time*

# Shortest Job First

| JOB | arrival_time<br>(s) | run_time<br>(s) |
|-----|---------------------|-----------------|
| A   | ~0                  | 60              |
| B   | ~0                  | 10              |
| C   | ~0                  | 10              |

What is the average turnaround time with SJF?

# SJF Turnaround Time



What is the average turnaround time with SJF?

$$(80 + 10 + 20) / 3 = \sim 36.7s$$

Average turnaround with FIFO: 70s

For minimizing average turnaround time (with no preemption):  
SJF is provably optimal

Moving shorter job before longer job improves turnaround time of short job more than it harms turnaround time of long job

# Scheduling Basics

---

## Workloads:

arrival\_time

run\_time

## Schedulers:

FIFO

SJF

STCF

RR

## Metrics:

turnaround\_time

response\_time

# Workload Assumptions

---

- ~~1. Each job runs for the same amount of time~~
- ~~2. All jobs arrive at the same time~~
3. All jobs only use the CPU (no I/O)
4. The run-time of each job is known

# Shortest Job First (Arrival Time)


---

| JOB | arrival_time<br>(s) | run_time<br>(s) |
|-----|---------------------|-----------------|
| A   | ~0                  | 60              |
| B   | ~10                 | 10              |
| C   | ~10                 | 10              |

What is the average turnaround time with SJF?



# Stuck Behind a Tractor Again



| JOB | arrival_time<br>(s) | run_time<br>(s) |
|-----|---------------------|-----------------|
| A   | ~0                  | 60              |
| B   | ~10                 | 10              |
| C   | ~10                 | 10              |

What is the average turnaround time?

$$(60 + (70 - 10) + (80 - 10)) / 3 = \mathbf{63.3s}$$

# Preemptive Scheduling

---

## Prev schedulers:

- FIFO and SJF are non-preemptive
- Only schedule new job when previous job voluntarily relinquishes CPU (performs I/O or exits)


## New scheduler:

- Preemptive: Potentially schedule different job at any point by taking CPU away from running job
- STCF (Shortest Time-to-Completion First)
- Always run job that will complete the quickest

# NON-PREEMPTIVE: SJF

| JOB | arrival_time<br>(s) | run_time<br>(s) |
|-----|---------------------|-----------------|
|-----|---------------------|-----------------|

|   |     |    |
|---|-----|----|
| A | ~0  | 60 |
| B | ~10 | 10 |
| C | ~10 | 10 |




Average turnaround time:

$$(60 + (70 - 10) + (80 - 10)) / 3 = \mathbf{63.3s}$$

# PREEMPTIVE: STCF

| JOB | arrival_time<br>(s) | run_time<br>(s) |        |
|-----|---------------------|-----------------|--------|
| A   | ~0                  | 60              | A: 80s |
| B   | ~10                 | 10              | B: 10s |
| C   | ~10                 | 10              | C: 20s |



Average turnaround time with STCF?

**36.6**

Average turnaround time with SJF: **63.3s**

# Scheduling Basics

---

## Workloads:

arrival\_time  
run\_time

## Schedulers:

FIFO  
SJF  
STCF  
RR

## Metrics:

turnaround\_time  
response\_time

# Response Time

---

Sometimes care about when job starts instead of when it finishes

New metric:


$$\text{response\_time} = \text{first\_run\_time} - \text{arrival\_time}$$



# Response vs. Turnaround

B's turnaround: 20s  $\longleftrightarrow$

B's response: 10s  $\longleftrightarrow$



# Round-Robin Scheduler

---


## **Prev schedulers:**

FIFO, SJF, and STCF can have poor response time

## **New scheduler: RR (Round Robin)**


Alternate ready processes every fixed-length time-slice

# FIFO vs RR



Avg Response Time?

$$(0+5+10)/3 = 5$$



Avg Response Time?

$$(0+1+2)/3 = 1$$

In what way is RR worse?

Ave. turn-around time with equal job lengths is horrible

Other reasons why RR could be better?

If don't know run-time of each job, gives short jobs a chance to run and finish fast

# Scheduling Basics

---

## Workloads:

arrival\_time  
run\_time

## Schedulers:

FIFO  
SJF  
STCF  
RR

## Metrics:


turnaround\_time  
response\_time

# Workload Assumptions

---


- ~~1. Each job runs for the same amount of time~~
- ~~2. All jobs arrive at the same time~~
- ~~3. All jobs only use the CPU (no I/O)~~
4. The run-time of each job is known

# Not I/O Aware



Don't let Job A hold on to CPU while blocked waiting for disk

# I/O Aware (Overlap)



Treat Job A as 3 separate CPU bursts

When Job A completes I/O, another Job A is ready

Each CPU burst is shorter than Job B, so with SCTF,  
Job A preempts Job B



# Workload Assumptions

---

- ~~1. Each job runs for the same amount of time~~
- ~~2. All jobs arrive at the same time~~
- ~~3. All jobs only use the CPU (no I/O)~~
- ~~4. The run-time of each job is known~~  
(need smarter, fancier scheduler)

## (Multi-Level Feedback Queue)

---

Goal: general-purpose scheduling

Must support two job types with distinct goals

- “**interactive**” programs care about **response time**
- “**batch**” programs care about **turnaround time**


Approach: multiple levels of round-robin;


each level has higher priority than lower levels and preempts them

# Priorities

Rule 1: If  $\text{priority}(A) > \text{Priority}(B)$ , A runs

Rule 2: If  $\text{priority}(A) == \text{Priority}(B)$ , A & B run in RR

Q3 → 

Q2 → 

Q1

Q0 →  → 

“Multi-level”

How to know how to set  
priority?

Approach 1: nice

Approach 2: history  
“feedback”

# History

---

- ② Use past behavior of process to predict future behavior
  - Common technique in systems
- ② Processes alternate between I/O and CPU work
- ② Guess how CPU burst (job) will behave based on past CPU bursts (jobs) of this process

# More MLFQ Rules

---

Rule 1: If  $\text{priority}(A) > \text{Priority}(B)$ , A runs

Rule 2: If  $\text{priority}(A) == \text{Priority}(B)$ , A & B run in RR

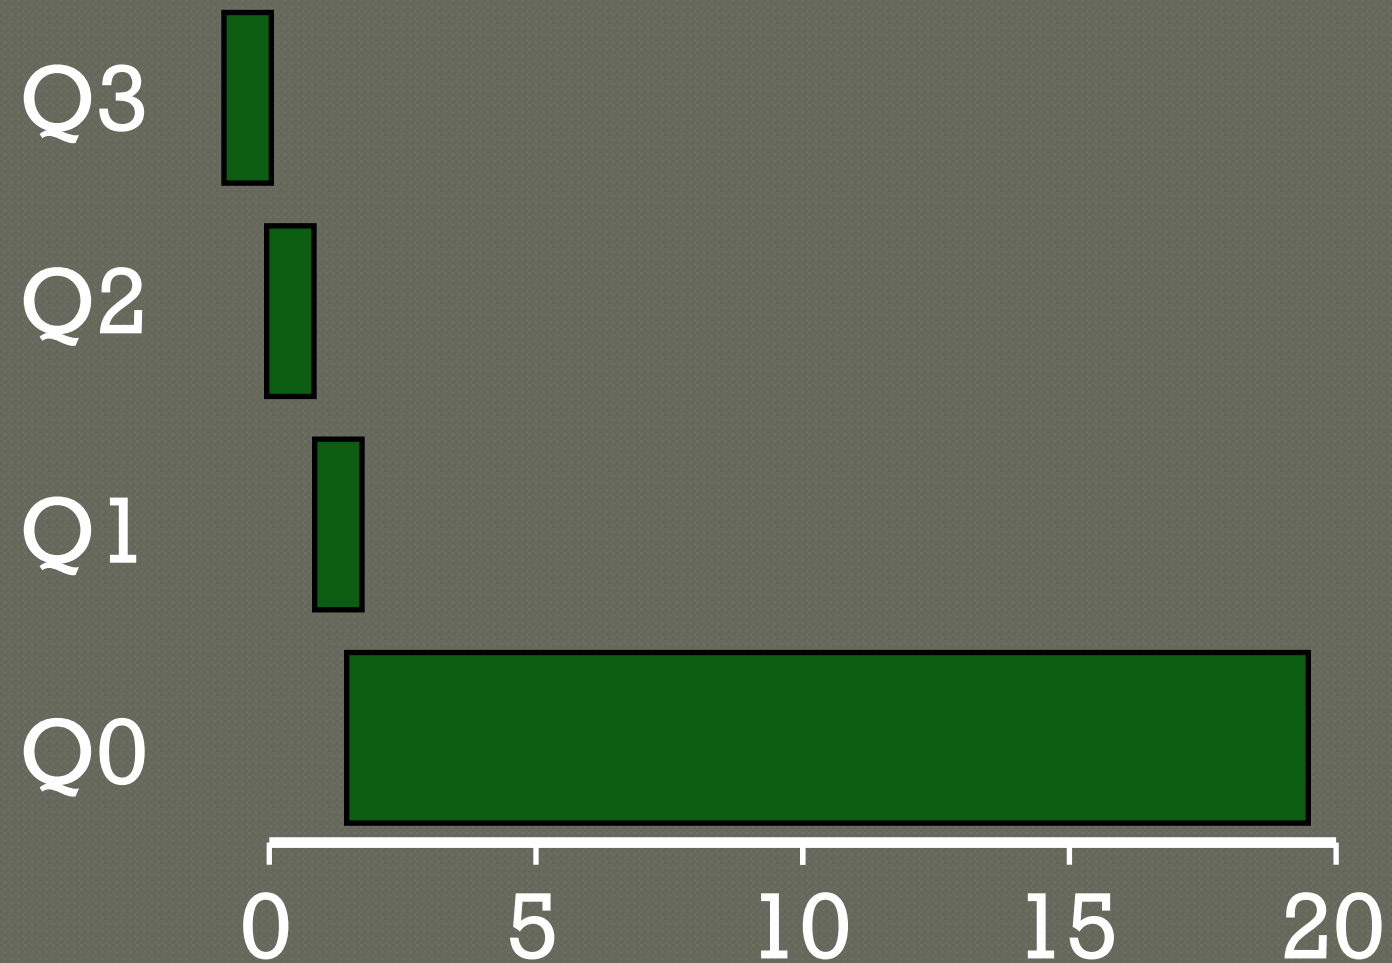
More rules:

Rule 3: Processes start at top priority


Rule 4: If job uses whole slice, demote process  
(longer time slices at lower priorities)

# One Long Job (Example)

---




# An Interactive Process Joins



Interactive process never uses entire time slice, so never demonstrates its priority




# Problems with MLFQ?



## Problems

- unforgiving + starvation
- gaming the system


# Prevent Starvation



Problem: Low priority job may never get scheduled

Periodically boost priority of all jobs (or all jobs that haven't been scheduled)

# Prevent Gaming



**Problem:** High priority job could trick scheduler and get more CPU by performing I/O right before time-slice ends

**Fix:** Account for job's total run time at priority level  
(instead of just this time slice);  
downgrade when exceed threshold

# Lottery Scheduling

---

Goal: proportional (fair) share

Approach:

- give processes lottery tickets
- whoever wins runs
- higher priority => more tickets

Amazingly simple to implement

# Lottery Code

---

```
int counter = 0;
int winner = getrandom(0, totaltickets);
node_t *current = head;
while (current) {
    counter += current->tickets;
    if (counter > winner) break;
    current = current->next;
}
// current is the winner
```

# Lottery example

```
int counter = 0;
int winner = getRandom(0, totaltickets);
node_t *current = head;
while(current) {
    counter += current->tickets;
    if (counter > winner) break;
    current = current->next;
}
// current gets to run
```

Who runs if **winner** is:

50

350

0



# Other Lottery Ideas

---

Ticket Transfers

Ticket Currencies

Ticket Inflation

(read more in OSTEP)



# Summary

---

Understand goals (metrics) and workload, then design scheduler around that

General purpose schedulers need to support processes with different goals

Past behavior is good predictor of future behavior

Random algorithms (lottery scheduling) can be simple to implement, and avoid corner cases.