


other methods as well

Hashed (for OSs > 32 bits)

Inverted

combined Paging & Segmentation

Virtual Memory



on 3.25 GHz machine
reg = 3.1×10^{-10}
cache = 1.5×10^{-10}
mem = 1×10^{-7}
disk = 1×10^{-3}

reg is 3000 times faster than mem

Motivation

Previously:

All of a process must be in mem
large process \rightarrow out of mem :/

But remember:

locality of ref: you tend to access mem around where you last accessed mem.

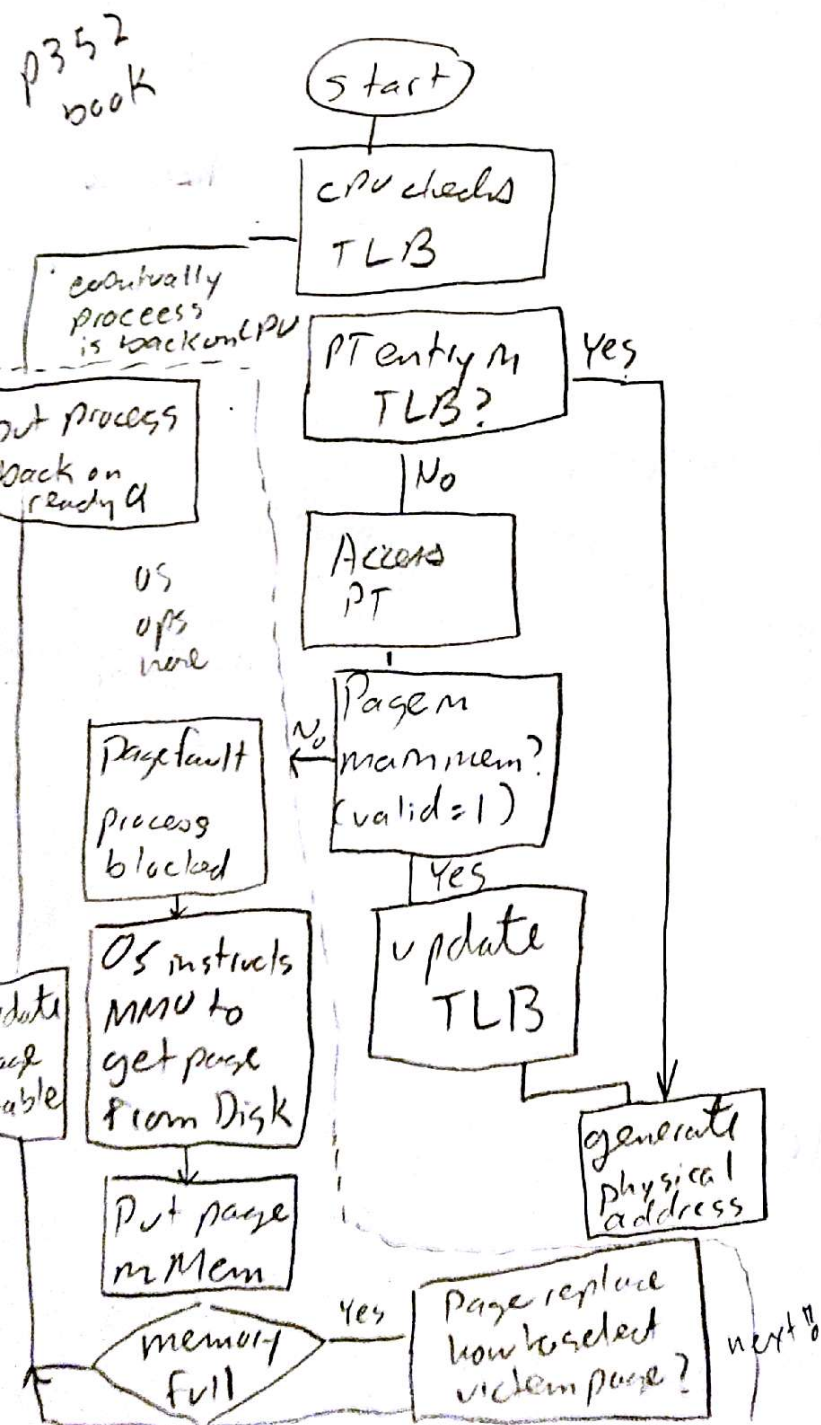
So...

Only need a small part of a process in mem at a time (part that's running)

Idea

Process runs when not all pages in memory

- keep close by pages in mem, those currently referenced (PC pointing to, function just called...)
- keep unreferenced stuff on HD
- Pull in from Disk when needed (X parent to the user)



Disk is 10,000 → 100,000 times slower than mem.
Big cost so load big pages when you have to go to disk (4K → 1Meg)

Virtual Mem operations.

Detect page fault

choose free physical page

OS algo. if none free OS chooses victim page (from current process)

Bring page to mem from disk


Need HW + Software

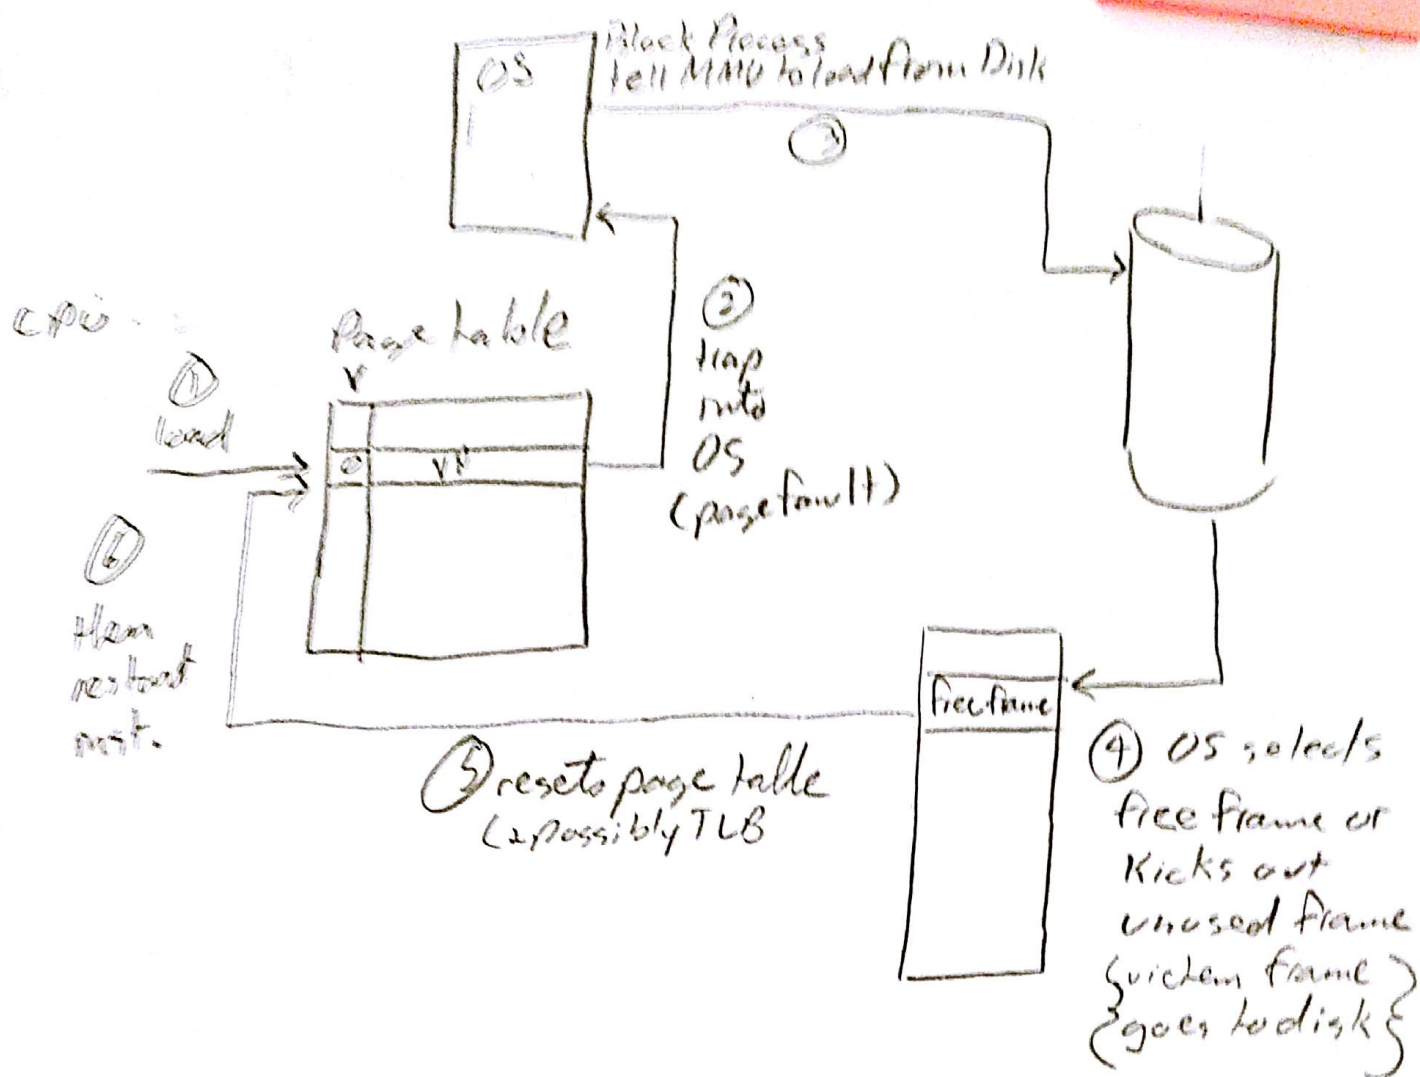
→ valid bit - if 0 then page fault.

$V=0$ if page on disk, store disk location in page table

$V=0$ page fault, referencing results in a trap to the OS

In OS page fault handler, check to see if fault is caused by ref to true invalid page or page on disk





when to bring pages from disk?

on Demand? Start up with no pages loaded
Wait until a page must be in memory

Request - user specifies, user manages memory by hand :C, User not expert, user can use up all memory

Prepaging - load page before needed, JIT, very efficient, when one page created bring in next, or predict what is needed (hard to do).