

**Department of Physics,  
Computer Science & Engineering**

CPSC 410 – Operating Systems I

# Chapter 2: Operating System Overview

**Keith Perkins**

Original slides by Dr. Roberto A. Flores

# Chapter 2 Topics

---

## OS functions

- Objectives, OS as user/computer interface, OS as resource manager

## OS evolution

- Serial, Batch, Multi-programming, Time sharing

## Achievements

- Process, Memory management, Information security, Scheduling, System structure

## Virtual machines

- Virtualization, Architecture

# OS Functions

---

## ● Functions

- User/Computer interface
  - An interface between applications and hardware
- Resource manager
  - A program controlling execution of application programs

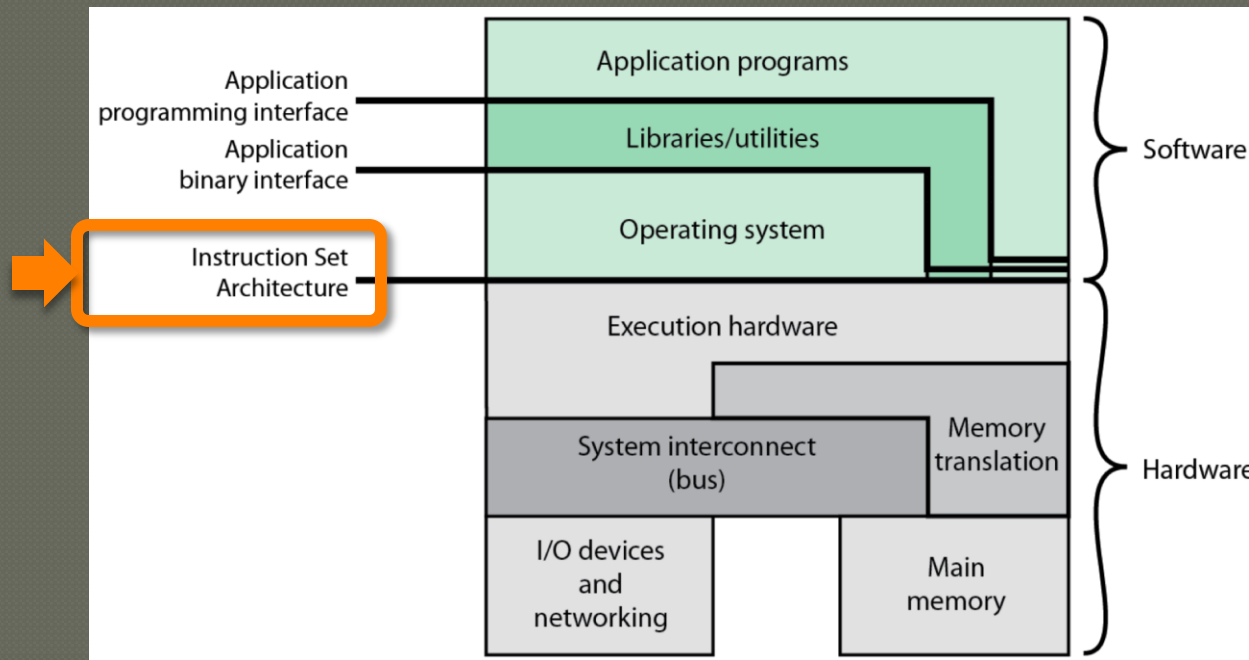
## ● Objectives

- Convenience
- Efficiency
- Ability to evolve

# User/Computer Interface

## Key interfaces

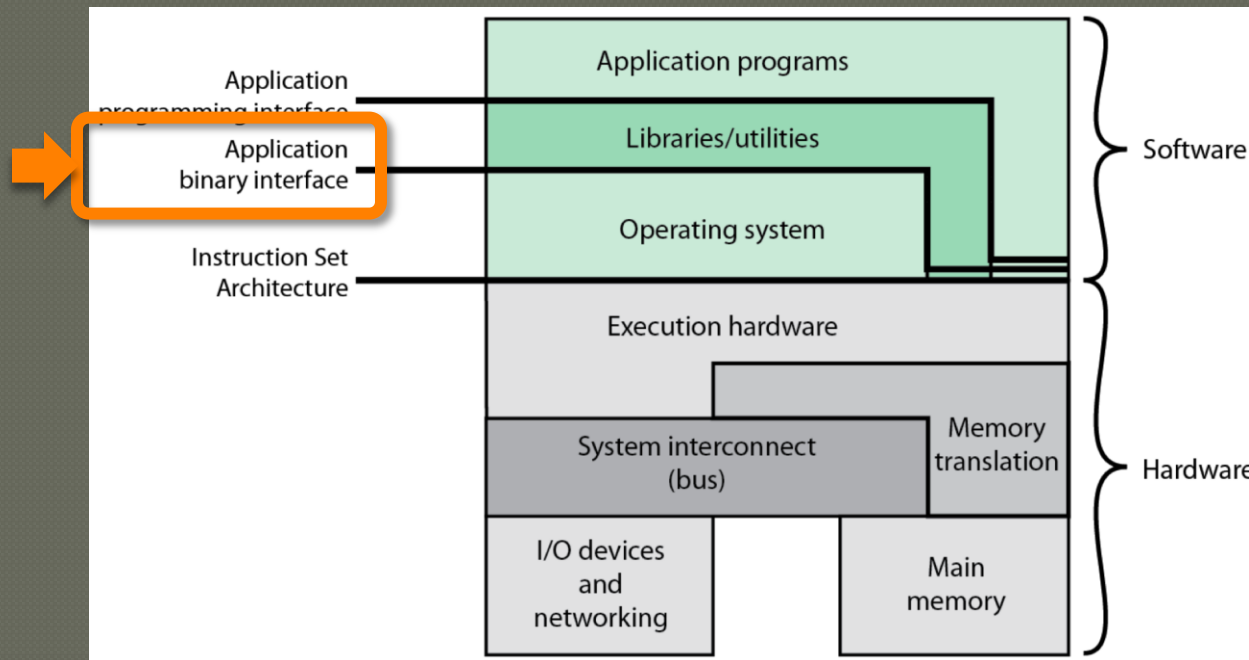
- **ISA** : Instruction set architecture
  - Machine language instructions hardware can execute
    - add registers, fetch memory, ...



# User/Computer Interface

## Key interfaces

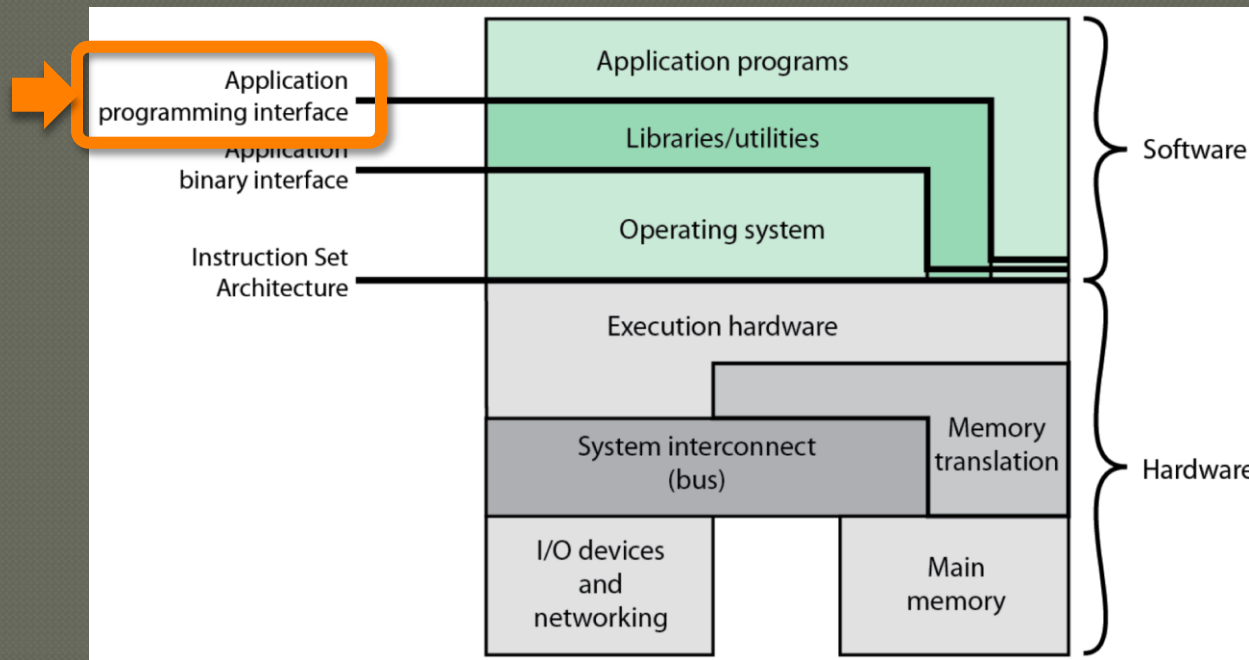
- **ABI** : Application binary interface
  - Masks hardware details
  - Mediate between programs & computer resources/services



# User/Computer Interface

## Key interfaces

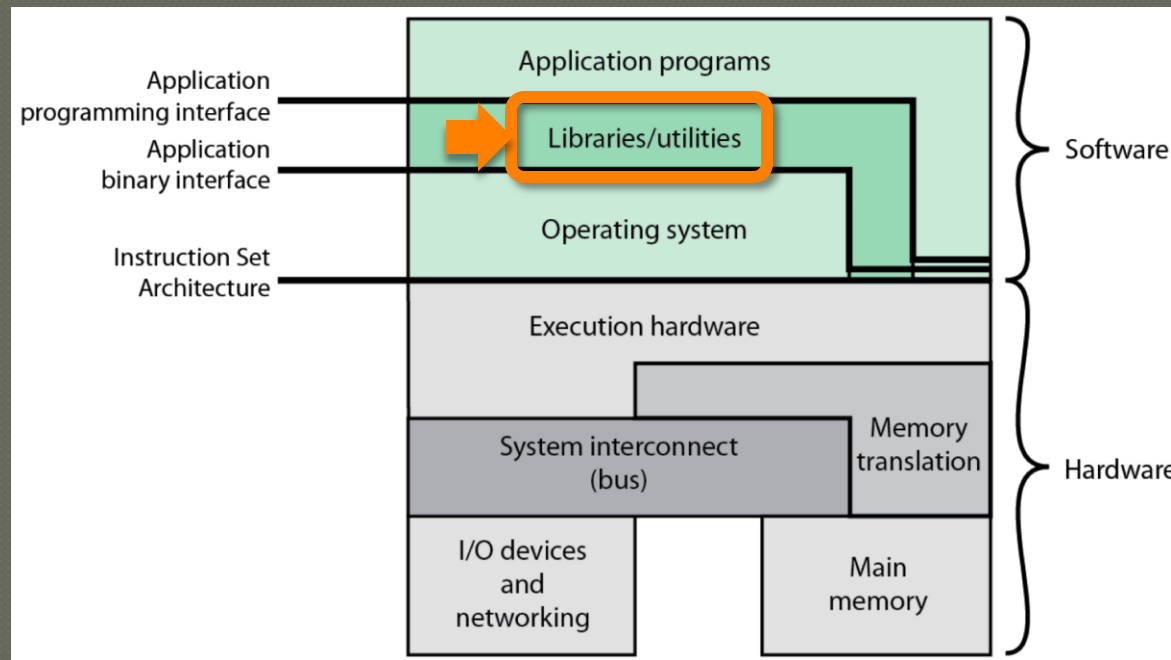
- **API** : Application programming interface
  - High-level language instructions
  - Facilitates source code portability (re-compilation)



# Resource Manager

## ● OS comes with libraries

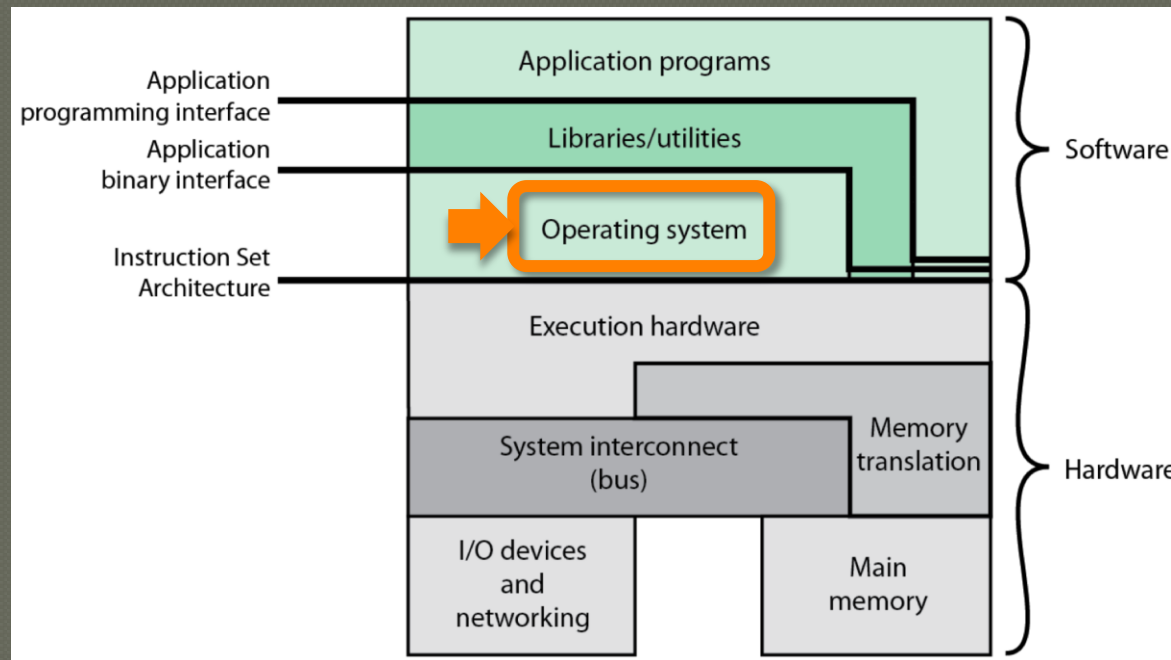
- Implementing functions to support creating programs, managing files, and controlling I/O devices
  - editors, interpreters, I/O modules, ...



# Resource Manager

## ● OS

- Controls hardware resources
- Manages system **services**
  - printer spooling, audio, sockets, ...





# Resource Manager

What is the OS doing for me?

## ● OS services

- Program development
  - editors, compilers, debuggers (not OS)
- Program execution
  - load data & instructions into memory
  - initialize I/O devices & files
- Access to I/O devices
  - uniform interface to access I/O devices (read, write)
    - Standard API (like read/write) to vendor supplied device drivers
- Controlled access to files
  - control orderly access to files (data integrity).

# Resource Manager

What is the OS doing for me?

## ● OS services (II)

- System access
  - control access to resources (permissions)
  - resolve conflicts for resource contention (privileges).
- Error detection & response
  - minimize disrupting running programs (least disruptive)
    - reporting error, retrying operation, ending erring program
- Accounting
  - Keeping logs & statistics

# Chapter 2 Topics

---

## ● OS functions

- Objectives, OS as user/computer interface, OS as resource manager

## ● OS evolution

- Serial, Batch, Multi-programming, Time sharing

## ● Achievements

- Process, Memory management, Information security, Scheduling, System structure

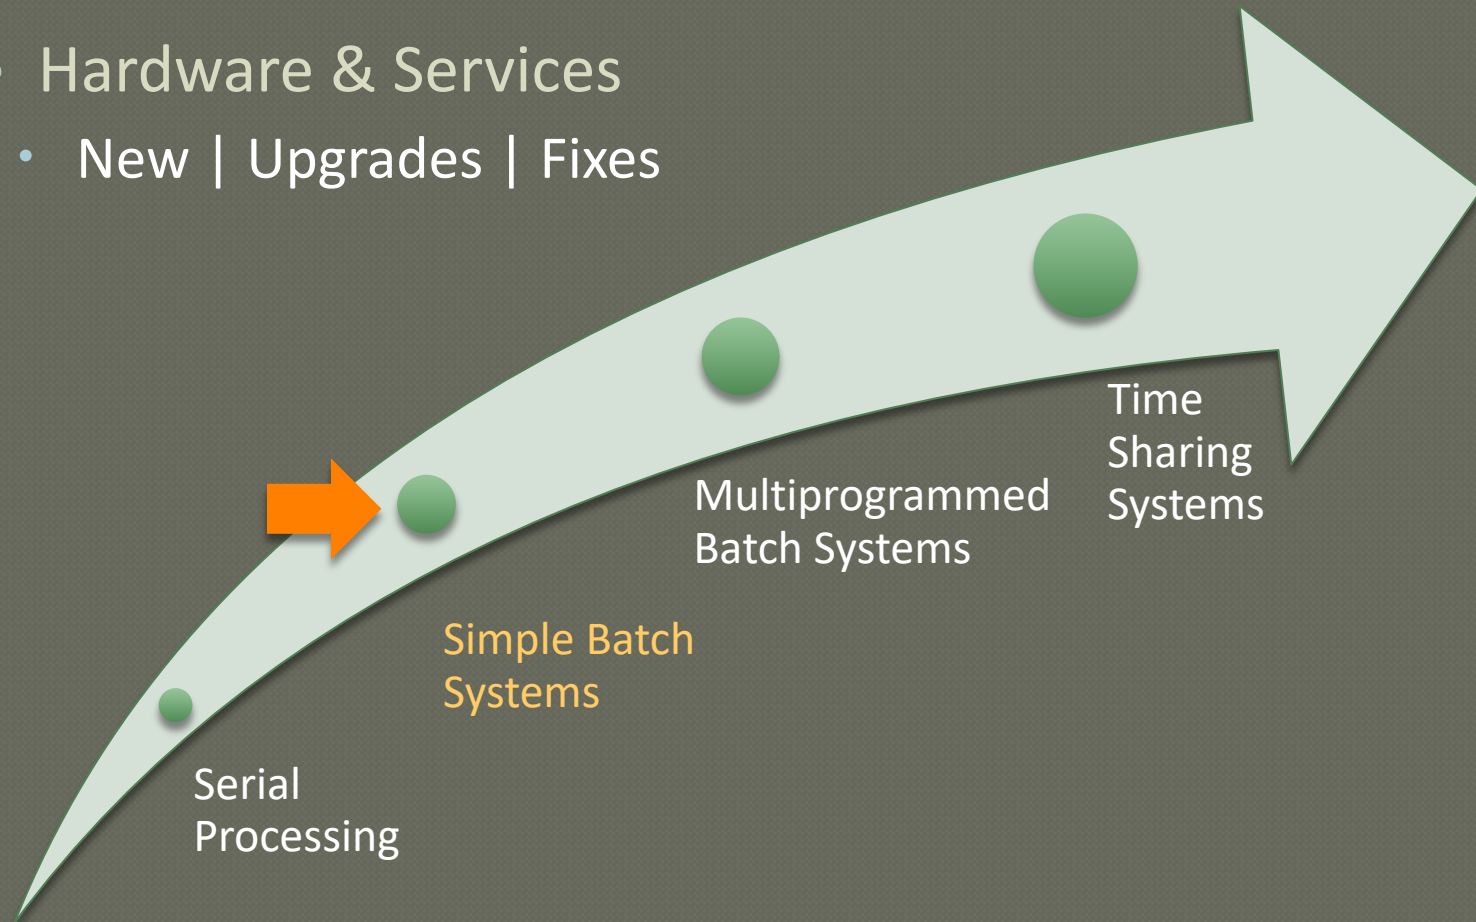
## ● Virtual machines

- Virtualization, Architecture

# Evolution

## Reasons for OS to evolve

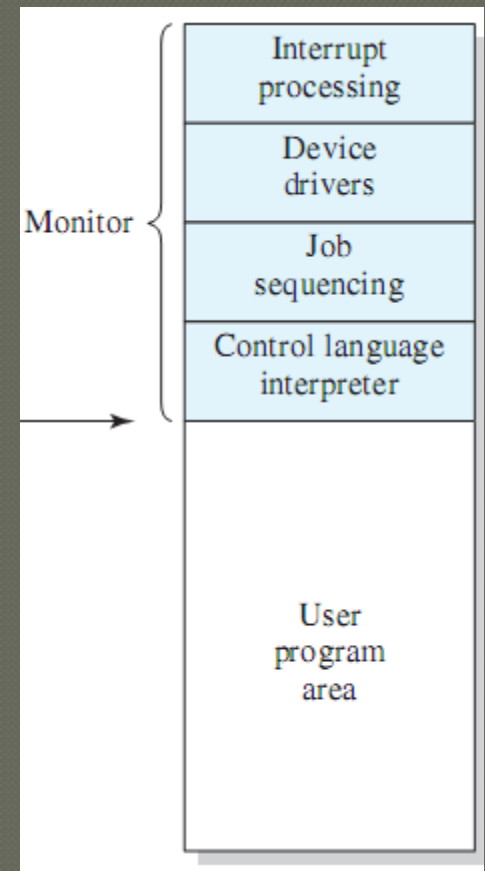
- Hardware & Services
  - New | Upgrades | Fixes



# OS Evolution

## Simple Batch Systems

- improving computer utilization
  - programmer has no direct access to computer
  - operator batches jobs, feeds them to an input device, then...
- Monitor (aka Batch OS)
  - program controlling the execution of jobs
  - 1. monitor reads next job & yields control of CPU to the job
    - “control is passed to a job” : CPU starts running user program
  - 2. user program ends & monitor continues running again
    - “control is returned to the monitor” : CPU runs monitor



# OS Evolution

---

## ● Simple Batch Systems (II)

- Job Control Language (JCL)
  - Instructions meant for the monitor (like pre-processing)
    - \$JOB \$FTN <source code> \$LOAD \$RUN <data> \$END
- Hardware support of Monitor
  - Memory protection
    - Memory where monitor resides is out-of-bounds for jobs
  - Timer
    - Notifies when jobs run longer than anticipated
  - Privileged instructions
    - Instructions that only the monitor can execute (e.g., load job)
  - Interrupts
    - Signals giving CPU a degree of flexibility

# OS Evolution

## Simple Batch Systems (II)

- Job Control Language (JCL)
  - Instructions meant for the monitor (like pre-processing)
    - \$JOB \$FTN <source code> \$LOAD \$RUN <data> \$END

- Hardware support of Monitor



### Memory protection

- Memory where monitor runs
- Timer
  - Notifies when jobs run



### Privileged instructions

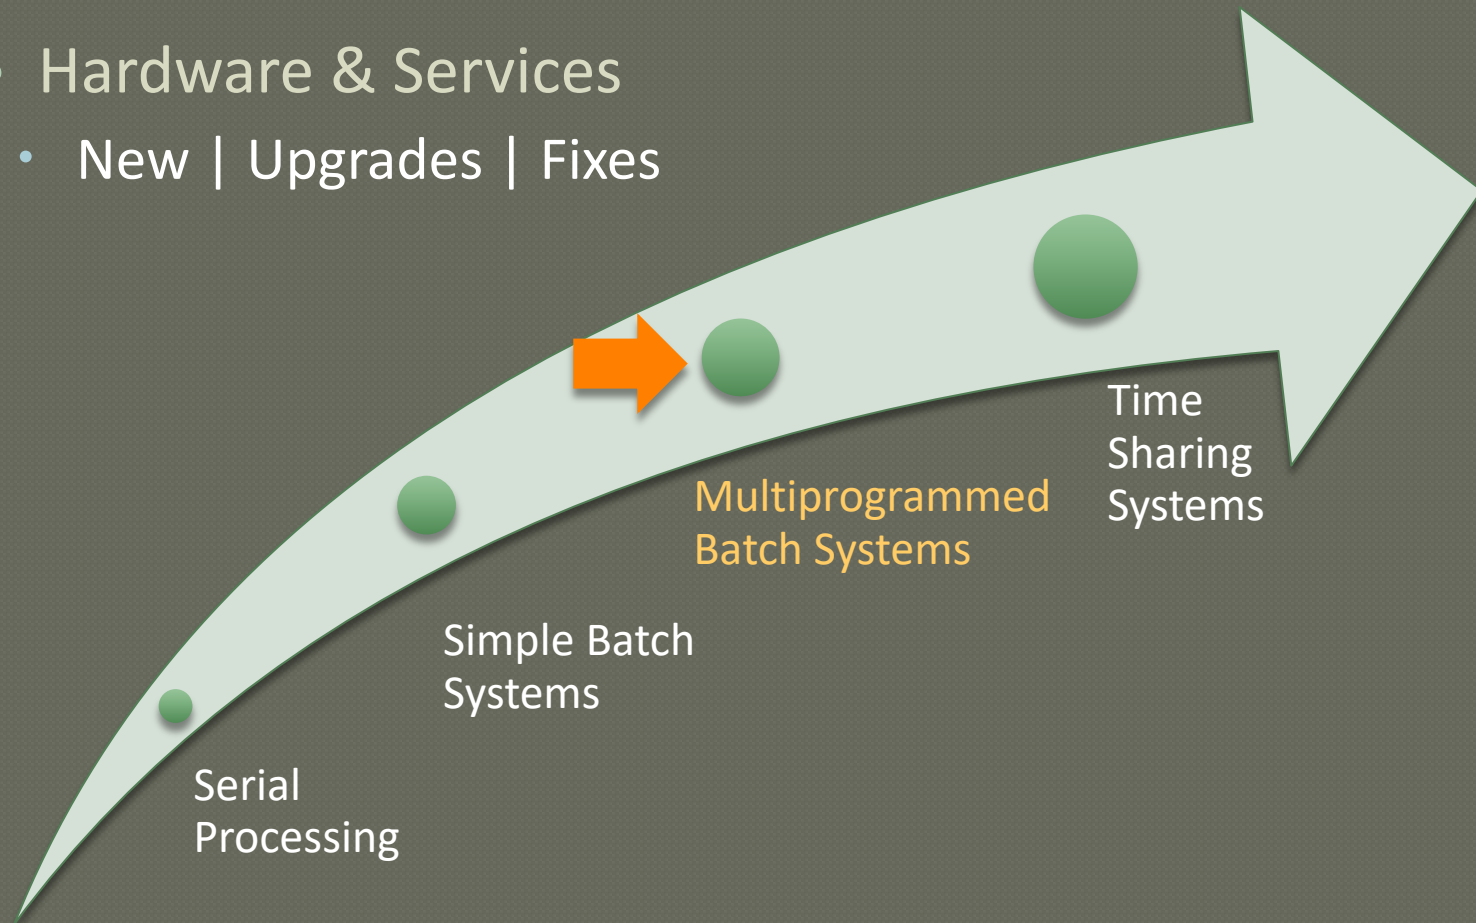
- Instructions that only the monitor can execute (e.g., load job)
- Interrupts
  - Signals giving CPU a degree of flexibility

	User Mode	Kernel Mode
Applies to...	User programs	Monitor
Memory access	Restricted	Unrestricted
Instructions	Limited	Unlimited

# Evolution

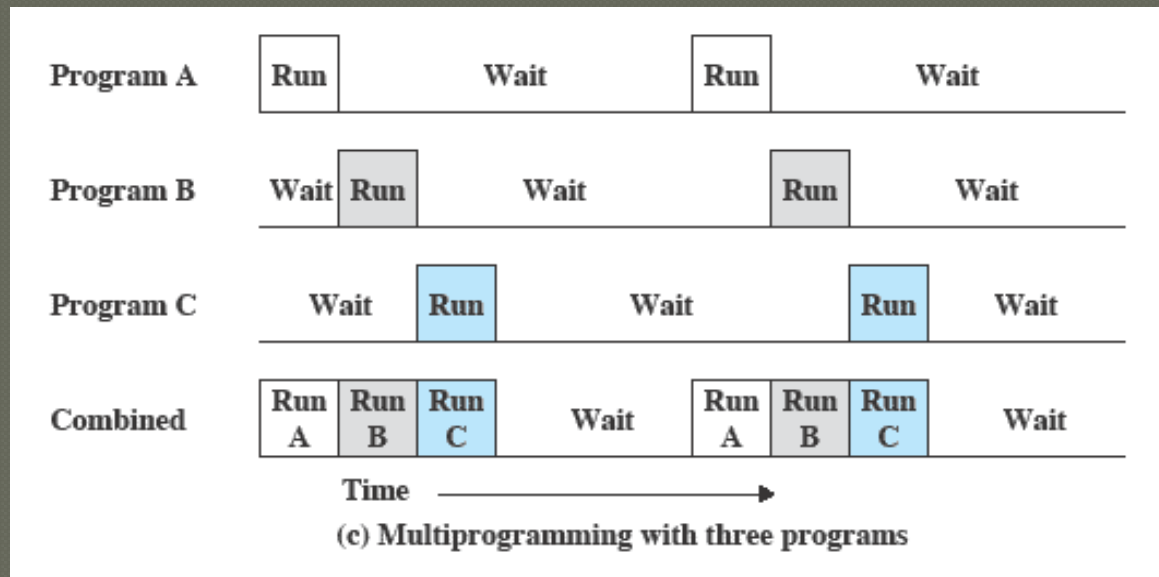
## Reasons for OS to evolve

- Hardware & Services
  - New | Upgrades | Fixes





# Multiprogramming



- Multiprogramming
  - also known as multitasking
  - memory is expanded to hold three, four, or more programs and switch among all of them

# Multiprogramming Example

**Table 2.1 Sample Program Execution Attributes**

	<b>JOB1</b>	<b>JOB2</b>	<b>JOB3</b>
<b>Type of job</b>	Heavy compute	Heavy I/O	Heavy I/O
<b>Duration</b>	5 min	15 min	10 min
<b>Memory required</b>	50 M	100 M	75 M
<b>Need disk?</b>	No	No	Yes
<b>Need terminal?</b>	No	Yes	No
<b>Need printer?</b>	No	No	Yes

# Utilization Histograms

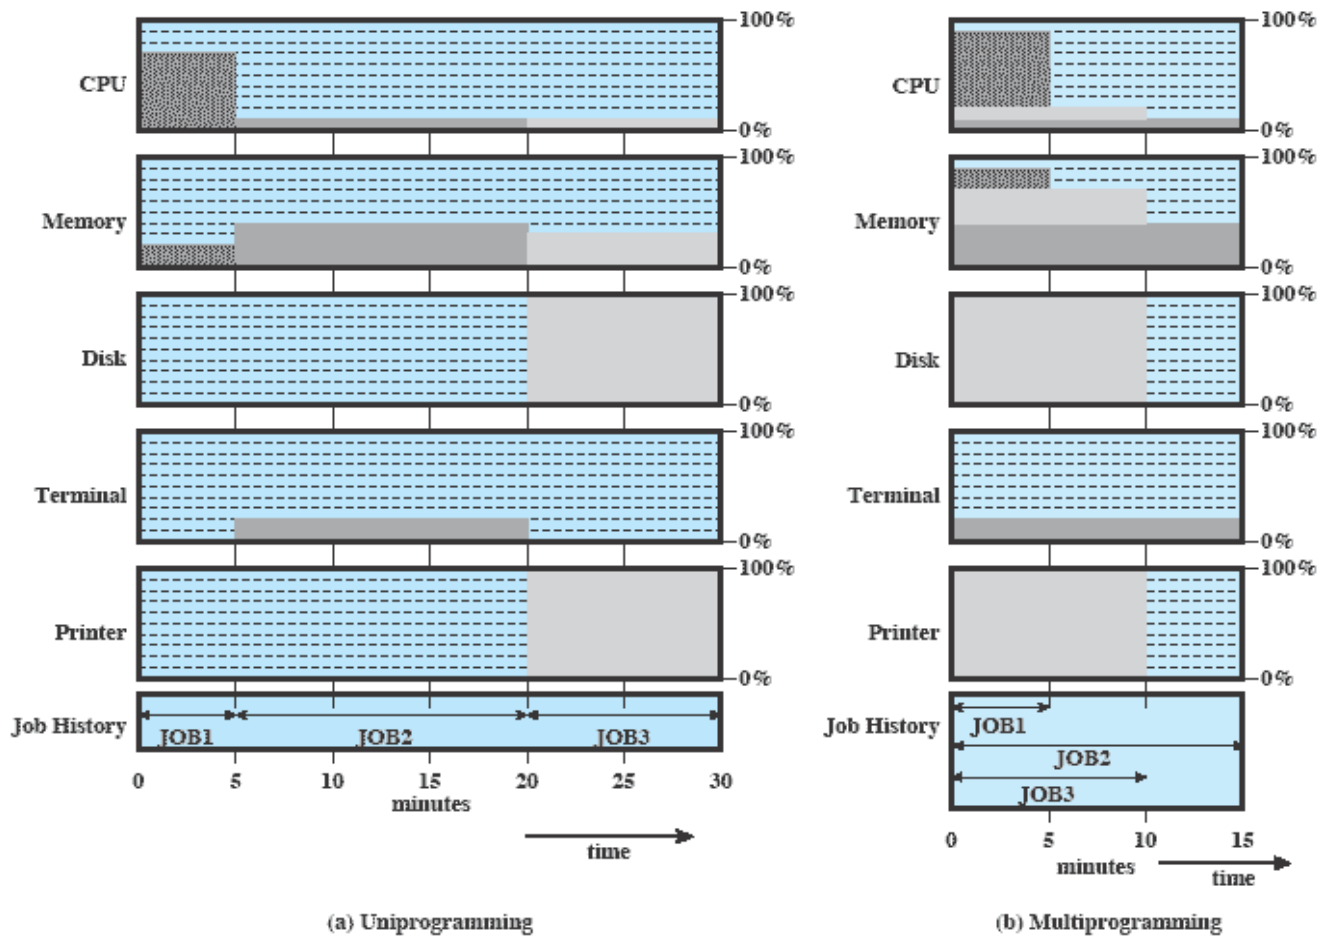


Figure 2.6 Utilization Histograms

# Effects on Resource Utilization

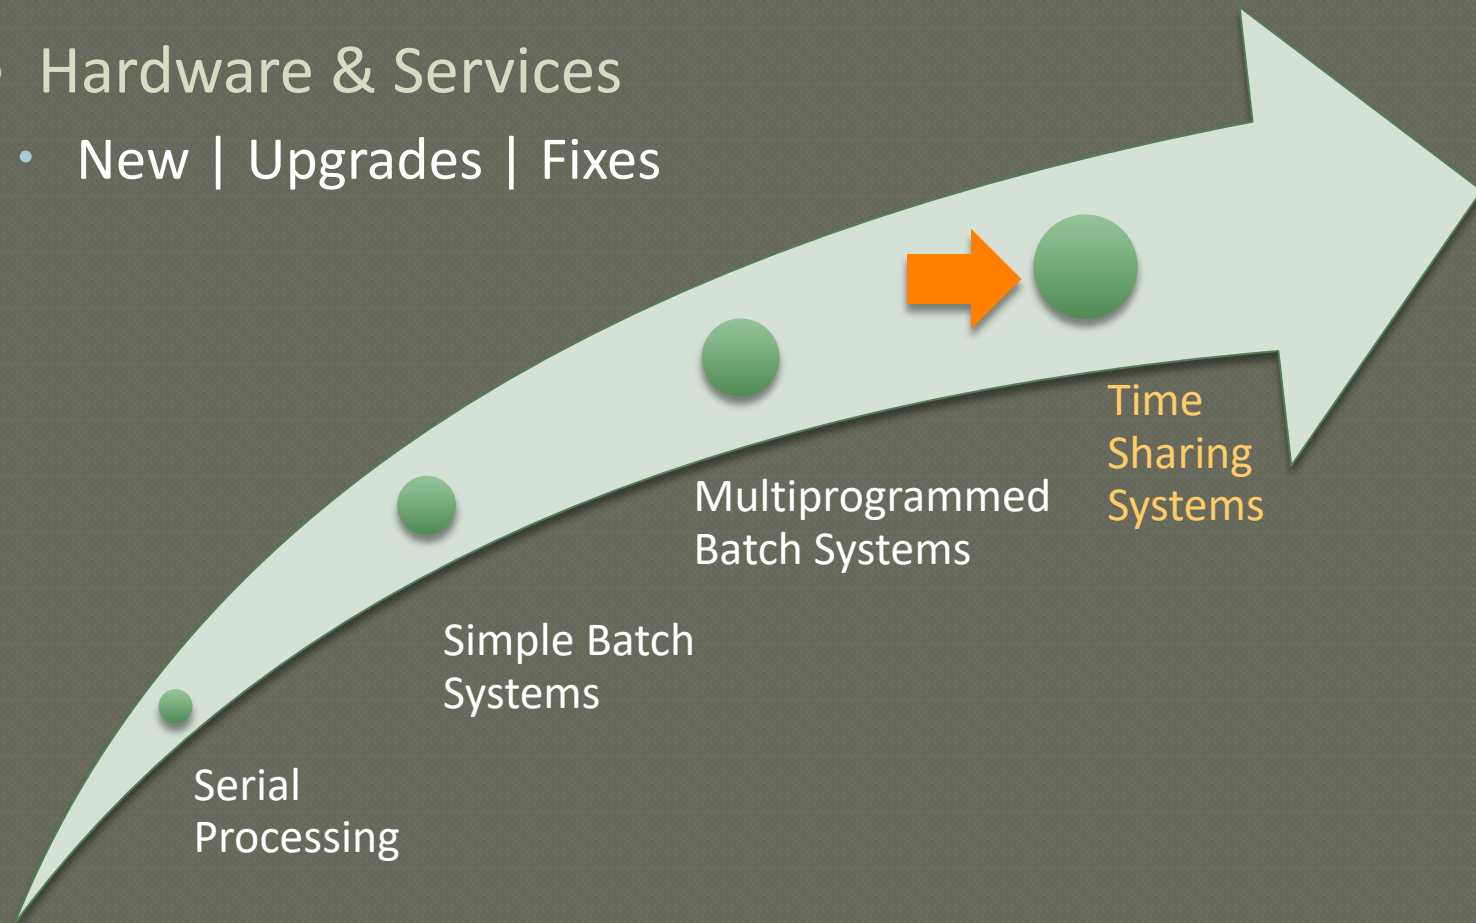
	Uniprogramming	Multiprogramming
<b>Processor use</b>	20%	40%
<b>Memory use</b>	33%	67%
<b>Disk use</b>	33%	67%
<b>Printer use</b>	33%	67%
<b>Elapsed time</b>	30 min	15 min
<b>Throughput</b>	6 jobs/hr	12 jobs/hr
<b>Mean response time</b>	18 min	10 min

Know how to calculate these numbers please

# Evolution

## Reasons for OS to evolve

- Hardware & Services
  - New | Upgrades | Fixes



# OS Evolution

## ● Time Sharing Systems

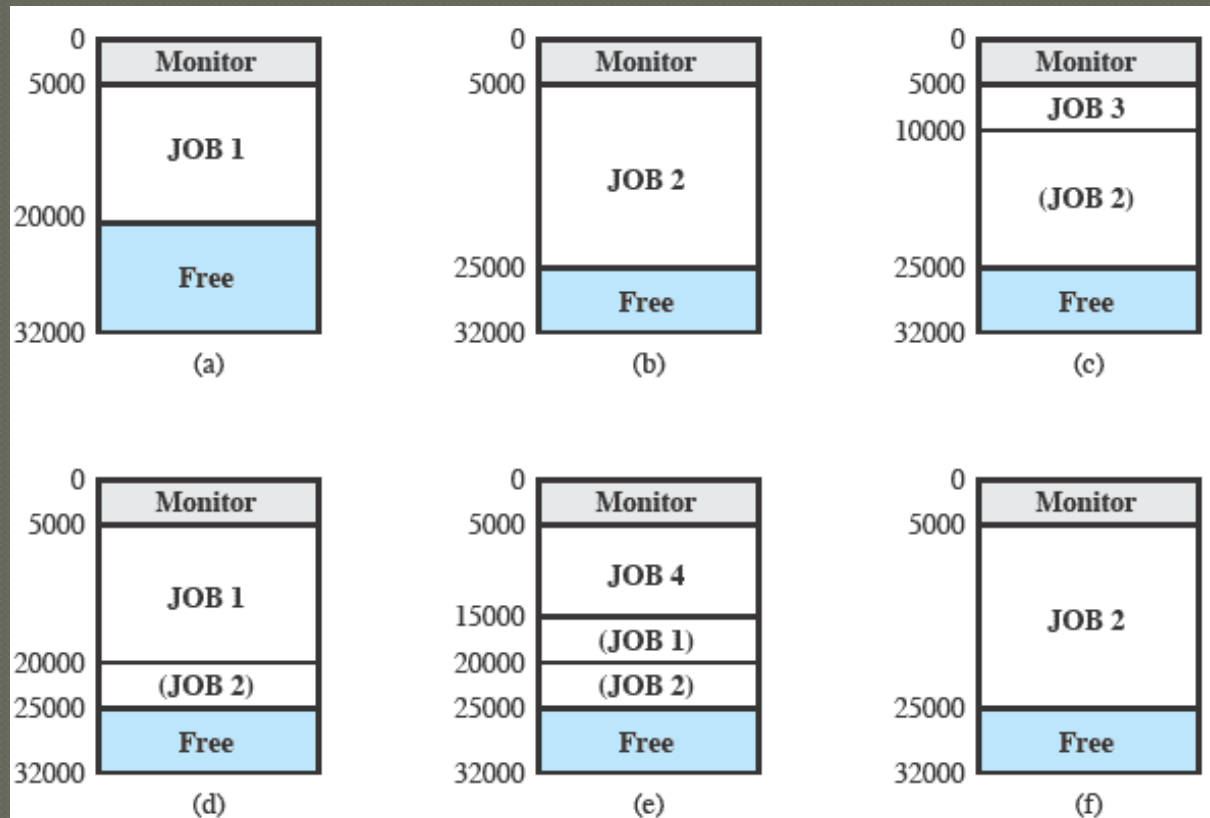
- Users access system simultaneously using terminals
- Time Slicing
  - Timer generates interrupts every 0.x seconds (small number)
  - OS preempts current program and loads in another
  - Preempted program & data are stored to disk (in old days)
  - *(but keep an eye on swapping overhead!)*
- Multi-Programming vs. Time sharing

	Batch Multi-programming	Time sharing
Objective	Maximize processor use	Minimize response time
Source of instructions	Job Control Language (JCL)	Commands entered in terminal

# OS Evolution

## Time Sharing Systems (II)

- CTSS: Compatible Time Sharing System
  - MIT (1961)
  - 32 users



# Chapter 2 Topics

---

## ● OS functions

- Objectives, OS as user/computer interface, OS as resource manager

## ● OS evolution

- Serial, Batch, Multi-programming, Time sharing

## ● Achievements

- Process, Memory management, Information security, Scheduling, System structure

## ● Virtual machines

- Virtualization, Architecture



# Achievements

---

## ● Major advances in OS development

- Processes
  - Definition, Errors, Components
- Memory management
  - OS responsibilities, Virtual memory
- Information protection & security
- Scheduling & resource management
- System structure

# Process

---

- Fundamental to the structure of operating systems

A *process* can be defined as:

a program in execution

an instance of a running program

the entity that can be assigned to, and executed on, a processor

a unit of activity characterized by a single sequential thread of execution, a current state, and an associated set of system resources

# Development of the Process

---

■ Three major lines of computer system development created problems in timing and synchronization that contributed to the development:

## **multiprogramming batch operation**

- processor is switched among the various programs residing in main memory with a goal of maximum efficiency

## **time sharing**

- be responsive to the individual user but be able to support many users simultaneously

## **real-time transaction systems**

- a number of users are entering queries or updates against a database

# Causes of Errors

## ● Improper synchronization

- a program must wait until the data are available in a buffer
- improper design of the signaling mechanism can result in loss or duplication



## ● Failed mutual exclusion

- more than one user or program attempts to make use of a shared resource at the same time

## ● Nondeterminate program operation

- program execution is interleaved by the processor when memory is shared
- the order in which programs are scheduled may affect their outcome

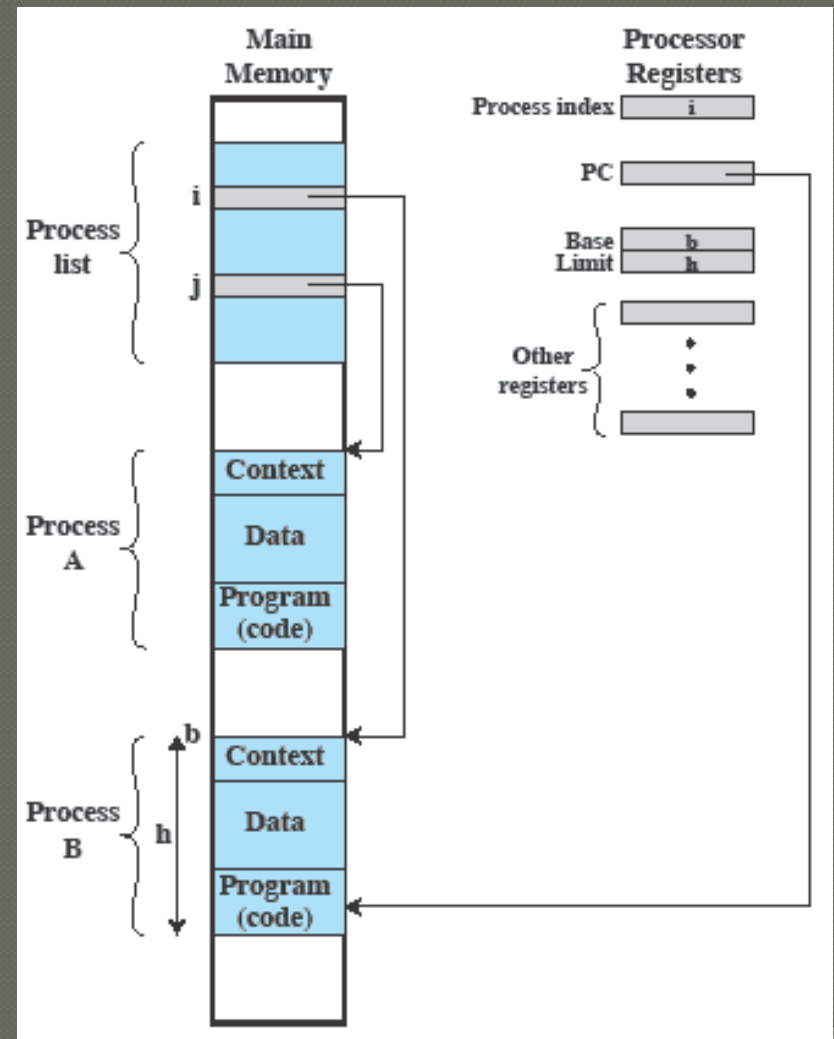
## ● Deadlocks

- it is possible for two or more programs to be hung up waiting for each other
- may depend on the chance timing of resource allocation and release

# Process Management

## Processes (components)

- Executable code
- Data
  - e.g., variables, buffers, ...
- Execution context (aka “process state”)
  - internal data used by the OS to control the process
    - e.g., registers, priority, whether it is waiting for an I/O event



# Achievements

## ● Memory management (OS responsibilities)

- Process isolation

Processes... are prevented from interfering with each other


- Automatic allocation & management
  - ...are not concerned about their own allocation
- Support of modular programming
  - ...are able to add/remove modules
- Protection & access control
  - ...are assured the integrity of data in shared memory
- Long-term storage
  - ...are able to store data for later runs (including power down)

How to handle simultaneous processes if they **do not fit** all in **main memory**?


# Achievements

- Memory management (Virtual Memory)

- Handling many processes with limited memory
- Paging
  - Processes are broken into blocks (aka **pages**)
    - Pages can be anywhere in main memory
  - CPU uses **virtual addresses** to find instructions/data
    - Addresses are **page** number + **offset** within page



# Virtual Memory Addressing



**Figure 2.10 Virtual Memory Addressing**



# Achievements

---

- Information protection & security
  - **controlling access** to processes & data
  - Availability
    - Protection against interruption
  - Confidentiality
    - Protection against unauthorized access
  - Data integrity
    - Protection against unauthorized modification
  - Authenticity
    - Protection against misrepresentation & data validation

# Achievements

---

## ● Scheduling & resource management

- OS **manages** resources (main memory, I/O devices, processors) and **schedules** their use by processes
- Fairness
  - Equal processes given equal and fair access to resources.
- Differential responsiveness
  - Different processes treated differently according to their needs.
- Efficiency
  - Overall performance is a goal
    - maximize throughput
    - minimize response time
    - accommodate as many users as possible

These criteria conflict (what's the right balance?)

# Key Elements of Schedule and Dispatch




Figure 2.11 Key Elements of an Operating System for Multiprogramming

# Achievements

## ● System structure

- Up to now

- OS are monolithic programs
- processes are linearly executed

What to do about it?

- **Microkernel Architecture**


- Keep essential functions in kernel
  - memory addressing, inter-process communication (IPC), scheduling
- Modularize the rest (towards **object-oriented** approach)
  - modules dynamically linked, easier to replace

- Advantages

- Flexibility: low coupling/high cohesion (props up **distributed OS** – illusion of unified memory & resources [in progress])

# Achievements

## Linux (flavor)



source: wikipedia.org

- Advantages
  - Flexibility: low coupling/high cohesion (props up **distributed OS** – illusion of unified memory & resources [in progress])


# Achievements

## ● System structure

- Up to now
  - OS are monolithic programs
  - processes are linearly executed
- **Symmetric multiprocessing** (add CPUs)
  - 2+ CPU run in parallel (hardware + OS exploiting it)
  - Processes scheduled to separate CPU (but share resources)
- **Multi-threading** (divide processes)
  - Process broken into parts that run concurrently (own thread)
  - Process =  $\sum$  (threads = concurrent unit of work)
  - Programmers control scope & timing of concurrency

What to do about it?

# Multitasking multiprogramming



**Figure 2.12** Multiprogramming and Multiprocessing

# Symmetric multiprocessing

## Challenges

- **Kernel concurrency:** Kernel processes allow concurrent CPU access (state integrity)
- **Scheduling:** Scheduling across CPUs must be coordinated (avoid duplicated runs)
- **Synchronization:** Access to resources must be synchronized (use locks)
- **Memory management:** Page reuse (coordinating page replacements)
- **Fault tolerance:** Graceful degradation

## Parallelism opportunities

- Multiprogramming & multi-threading in each processor
- A process could have its threads executed in different CPUs

Processes scheduled to separate CPU (but share resources)

- Multi-threading (divide processes)
  - Process broken into parts that run concurrently (own thread)
  - Process =  $\sum$  (threads = concurrent unit of work)
  - Programmers control scope & timing of concurrency



# Chapter 2 Topics

---

## ● OS functions

- Objectives, OS as user/computer interface, OS as resource manager

## ● OS evolution

- Serial, Batch, Multi-programming, Time sharing

## ● Achievements

- Process, Memory management, Information security, Scheduling, System structure


## ● Virtual machines

- Virtualization, Architecture

# Virtual Machines

## Virtualization


- [applies to OS | hardware]
- **when a computer runs** (simultaneously or not) **2+ OS** (different OS or different sessions of the same OS)
- OS can run 1+ **virtual machines** (VM)
  - Software implementing a particular OS or hardware
- Virtual Machine Monitor (VMM) (aka **hypervisor**)
  - Typical structure ➡



# Virtual Machines


## Process VM

- Each VM runs 1 process only
  - Process + VM start/stop as one
  - Portability through VM ABI
  - e.g., Java VM, MS .NET



## System VM

- Each VM runs 1+ processes
  - VM Monitor controls 1+ OS
  - Portability through VM API
  - **Hosted VM**: host has an OS
    - e.g., VMWare



# Chapter 2 Topics

---

## • OS functions

- Objectives, OS as user/computer interface, OS as resource manager

## • OS evolution

- Serial, Batch, Multi programming, Time sharing

## • Achievements

- Process, Memory management, Information security, Scheduling, System structure

## • Virtual machines

- Virtualization, Architecture

Done!