

# Mutual Exclusion

## Enforcement

only 1 thread in critical section of shared object.

## Availability

if no thread in critical section then any thread can enter.

## Minimal stay

Threads stay in critical section for minimal time

## Consistency

If resource must be protected, then every access to that resource is protected

(cannot have 1 access that is not protected) (remember)

---

## Hardware enforced.

what if we disable interrupts?

- guarantees atomic code

### Bad

- cannot have overlapping critical sections
- disables scheduling to other non-related (does not use shared resource) processes
- will not work multiprocessor (can't disable interrupts on 2 cores at a time.
- kills performance on single core

It will work on single core performance.

## Compare & Swap

atomic op, no race possible here

```
1 int cas ( int *word, int notlockedval, int lockedval)
2 {
3     int oldval = *word;
4     if (oldval == notlockedval) // if we can
5         *word = lockedval; // then do
6     return oldval;
7 }
```

lets say notlockedval = 0 } if \*word == 0 we can  
lockedval = 1 } lock it otherwise not.

const int NLV = 0;

const int LV = 1;

int word = NLV; // start off unlocked

void withdraw (int amount) {

stay in  
loop until  
cas() says  
NLV

```
1 while ( cas ( &word, NLV, LV) == LV) {
2     if (balance > amount) {
3         cout << "Approved";
4         balance -= amount;
5     }
6     word = NLV;
7 }
```

① 1<sup>st</sup> thrd → word = NLV; line 1 condition false goto 2

② interrupt at 3, 2<sup>nd</sup> thrd calls 1

③ cas returns LV == LV so it busy waits (wasting CPU time)

④ 1<sup>st</sup> thrd swapped back in & finishes, line 6 word = NLV

⑤ 2<sup>nd</sup> thrd line 1, cas returns NLV != LV so it proceeds.



4)

- bad



↳ very bad, watch CPU usage spike

→

- starvation & deadlock both possible

NO  $\rightarrow$  CAS must be atomic & my CAS is not  
 this function is ~~part of~~ the part of the HW

NO  $\rightarrow$  CAS must be atomic & my CAS is not  
~~this function is part of the part of the HW~~

It has to be part of hardware & atomic

ie (no interrupts)

## First

- talk about mutexes, C++ construct for synchronization between threads. (1 at a time)
  - with examples of how to solve withdrawal problem.
- then condition vars
- then semaphores, grouped up mutex, can allow more than 1 at a time
  - example
- then how semaphores implemented

mutex (thread based not process based) need shared memory.  
in linux, threads are treated as processes with the same mem space.

#include <mutex>

std::mutex g\_mutex;

void lock(); // if avail will proceed, otherwise locks

void unlock(); // unlocks (once per call);

bool trylock(); // locks if possible or returns false  
// no blocking

- do not call lock multiple times from same thread!  
trylock [if you ~~must~~ must use recursive-mutex]

- unlock mutex when you are done!

solve withdrawal probs

mutex g\_mutex;

void withdraw (int amount) {

g\_mutex.lock();

if (balance > amount) {

cout << "approved";

balance -= amount;

}

g\_mutex.unlock();

}

what happens if you throw an exception & never unlock  
it? (Deadlock) for all other threads Kill & restart.

better solution

how about a class?

.cpp

lockguard::lock\_guard (mutex &amutex) {  
g\_mutex = &amutex; amutex.lock();

lockguard::~lock\_guard () {  
(&amutex).unlock();

}

(auto unlocks when it goes  
out of scope)

class lock\_guard

{ private:

mutex\* g\_mutex;

public:

lock\_guard (mutex &amutex);

~lock\_guard();



ready in C++!

#include <mutex>

mutex gmutex;

lock\_guard<std::mutex> lock(gmutex);

↑  
when this goes out of scope it unlocks!

show in withdrawal problem.

show you do not have to unlock

but what if you want to lock  
across functions

```
int get balance() {
```

```
    return balance;
```

```
}
```

show separate function ~~with~~ call in withdraw (so  
its protected) is it OK? No! accessible ~~outside of lock~~ code

```
add balance (int i) {
```

```
    balance += i;
```

```
}
```

## Condition Variable

used to signal between threads. [thread(s)  
block until signaled)

overview

bool isReady = false; // global

T1 (waiting on isReady == true) ← blocked in meantime

T2 (sets isReady = true & then notifies any other threads waiting)

std::condition\_variable

### Notifying

notify\_one → os wakes 1 thread (can't tell which one)

notify\_all → os wakes them all (I will acquire mutex & continue) the rest wait until they get mutex & continue

### waiting

wait wait until notified

wait\_for } timed versions of  
wait\_until } above ignore for now



```
# include <mutex>
# include <condition-variable>
```

```
mutex m;
condition_variable c;
bool ready = false;
void get() {
```

```
    unique_lock<mutex> lock(m); ← acquire
    while (!ready)
        c.wait(lock) ← release & wait for signal
}
```

```
void set() {
```

```
    unique_lock<mutex> lock(m);
    ready = true; // all access to ready protected
    cv.notify_all; // wake up all other threads.
}
```

unique-lock verses  
lockguard?

lockguard is locked  
until destruct

unique-lock is  
all that can unlock,  
create unlocked then  
lock

unique-lock a little  
heavier, always use  
lockguard unless

you need  
unique lock  
abilities