



**Department of Physics,
Computer Science & Engineering**

CPSC 410 – Operating Systems I

Chapter 7: Memory Management

Memory Management

◉ Intro

◉ Requirements

- Relocation, Protection, Sharing, Logical & Physical organization

◉ Partitioning

- Fixed & Dynamic partitioning

◉ Paging

- Frames & pages, Addressing

● Memory Management

- **one part** of memory is used by the **OS**
- the **other** is used by **processes**
 - **Memory management** deals with the use and control of this memory among processes.

● Terminology

- **Frame** : a **fixed**-size block of main memory
- **Page** : a **fixed**-size block of virtual memory
- **Segment**: a **variable**-size block of...
...a **process** stored on **disk**

Memory Management

◉ Intro

◉ Requirements

- Relocation, Protection, Sharing, Logical & Physical organization

◉ Partitioning

- Fixed & Dynamic partitioning

◉ Paging

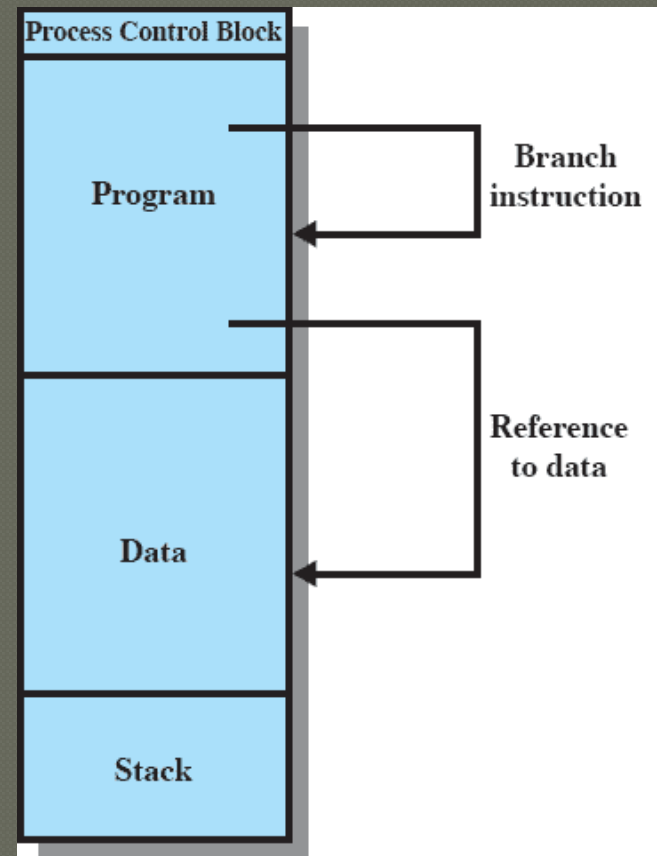
- Frames & pages, Addressing

Requirements

...that Memory Management is meant to satisfy

● Relocation

- processes are **loaded** to main memory to run.
- eventually, they are **swapped** in and out of main memory to maximize CPU utilization.
- **Relocation** implies that processes may get loaded into **different memory spaces** between swapping.
- This has implications for **addresses** within processes. ➡



Requirements

...that Memory Management is meant to satisfy

● Protection

- Are processes **referencing** correct memory locations?
 - locations may change between swaps
- ...memory references must be checked at runtime
 - **relocation** must also support **protection**

● Sharing

- Processes using the **same modules** could use one copy rather than having their own
 - **protection** must not be compromised when sharing memory
 - **relocation** must also support **sharing**

Requirements

...that Memory Management is meant to satisfy

◉ Logical organization

- Memory (main & secondary) are linear
- Programs are not! They use libraries (code abstraction)
 - written & compiled independently, can be shared

◉ Physical organization

- flow of information between main & secondary memory
 - loading/unloading modules & data
- Should programmers manage this flow?
 - What if a program + data does not fit into memory?
 - What if there are other programs running concurrently?
 - How much memory is available? Where/when will it become available?

This is getting too complicated!

Requirements

...that Memory Management is meant to satisfy

Logical organization

- Memory (main & secondary) are linear
- Programs are not! They use libraries (code abstraction)
 - written & compiled independently, can be shared

Physical organization

- flow of information between main & secondary memory
 - loading/unloading modules & data
- Should programmers manage this flow?
 - What if a program + data does not fit into memory?
 - What if there are other programs running concurrently?
 - How much memory is available? Where/when will it become

That's why memory management is needed

This is getting too complicated!

Memory Management

◉ Intro

◉ Requirements

- Relocation, Protection, Sharing, Logical & Physical organization

OK so...

memory management is all about

bringing processes into main memory for execution

- involves partitioning, paging & segmentation
 - (although **obsolete** they help contrasting other concepts)
- involves virtual memory
 - (currently **in use**, see next chapter)

Memory Management

- ◉ Intro

- ◉ Requirements

- Relocation, Protection, Sharing, Logical & Physical organization

- ◉ Partitioning

- Fixed & Dynamic partitioning

- ◉ Paging

- Frames & pages, Addressing

- ◉ Security issues

Partitioning

● Fixed Partitioning

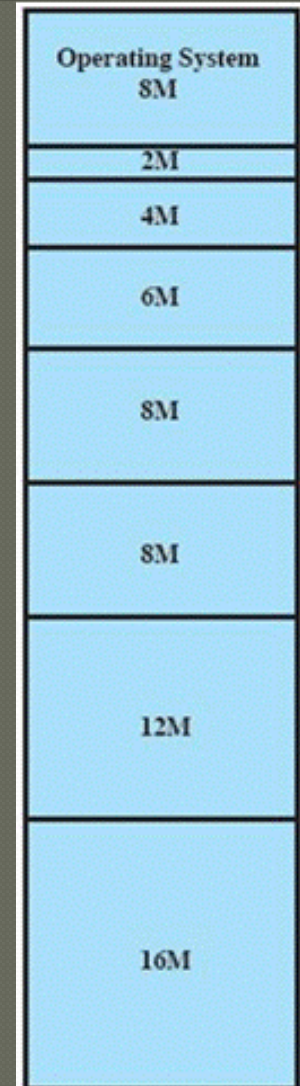
- Equal-size partitions
 - a process is loaded into a partition
 - OS swaps processes in & out as needed
- Disadvantages
 - What if a process is larger than a partition?
 - code must be designed with overlays
 - What if a process is smaller than a partition?
 - leftover memory is not used
 - aka internal fragmentation
 - wasted space due to the process loaded being smaller than the partition



Partitioning

● Fixed Partitioning (II)



- **Unequal-size** partitions
 - **Larger** processes can be accommodated without the need of overlays
 - There is **less** internal fragmentation by using best fit partition
- Disadvantages (ditto for Fixed)
 - Number of **partitions** (set at startup) limits the number of active processes
 - Prone to memory waste in cases when there are many **small jobs**



Partitioning

● Fixed Partitioning (II)


- Unequal-size partitions (placement algorithm)



Partitioning

● Dynamic Partitioning

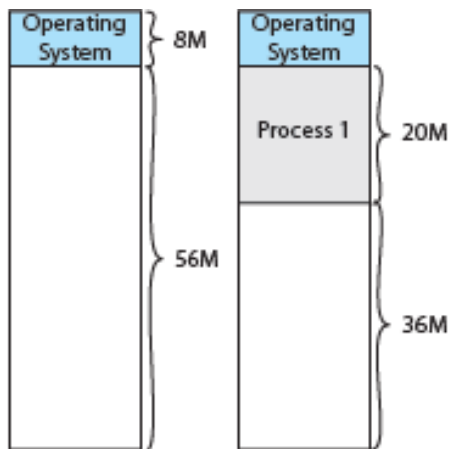
- Partitions vary in **length** & **number**
- Processes are given the **exact memory** they require
- Example (RAM 64M)
 - P1 starts (20M), P2 starts (14M), P3 starts (18M)
 - P2 ends, P4 starts (8M), P1 ends, P2 restarts (14M)



Partitioning

● Dynamic Partitioning

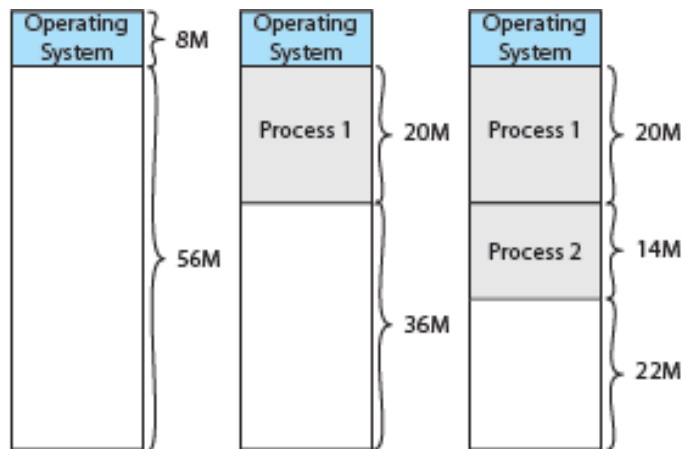
- Partitions vary in **length** & **number**
- Processes are given the **exact memory** they require
- Example (RAM 64M)
 - **P1 starts (20M)**, P2 starts (14M), P3 starts (18M)
 - P2 ends, P4 starts (8M), P1 ends, P2 restarts (14M)



Partitioning

● Dynamic Partitioning

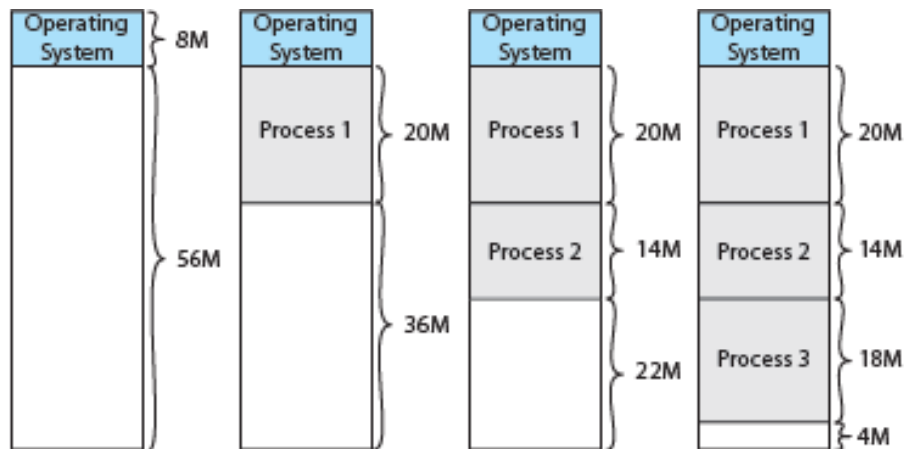
- Partitions vary in **length** & **number**
- Processes are given the **exact memory** they require
- Example (RAM 64M)
 - P1 starts (20M), **P2 starts (14M)**, P3 starts (18M)
 - P2 ends, P4 starts (8M), P1 ends, P2 restarts (14M)



Partitioning

● Dynamic Partitioning

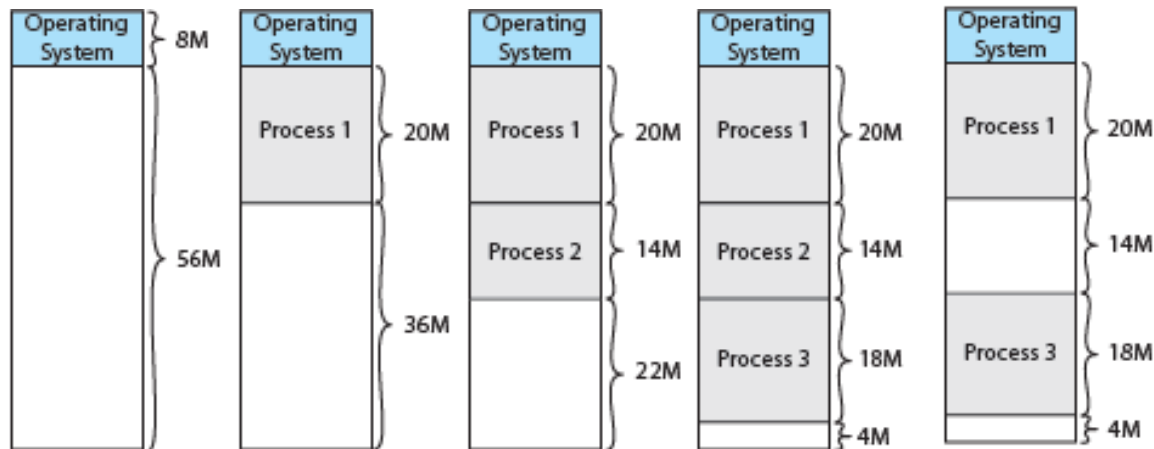
- Partitions vary in **length** & **number**
- Processes are given the **exact memory** they require
- Example (RAM 64M)
 - P1 starts (20M), P2 starts (14M), **P3 starts (18M)**
 - P2 ends, P4 starts (8M), P1 ends, P2 restarts (14M)



Partitioning

● Dynamic Partitioning

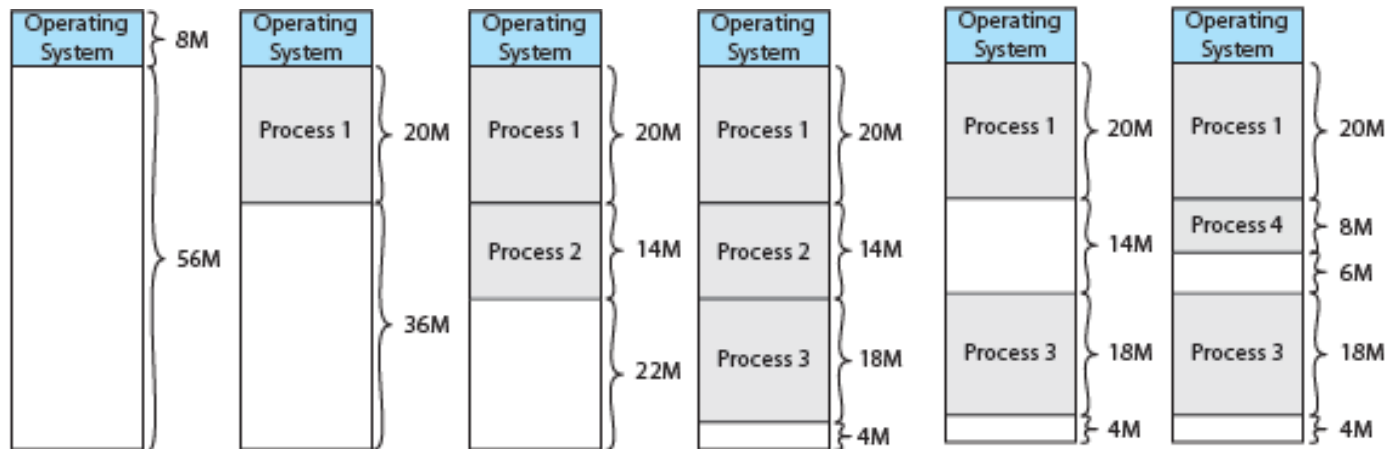
- Partitions vary in **length** & **number**
- Processes are given the **exact memory** they require
- Example (RAM 64M)
 - P1 starts (20M), P2 starts (14M), P3 starts (18M)
 - **P2 ends**, P4 starts (8M), P1 ends, P2 restarts (14M)



Partitioning

● Dynamic Partitioning

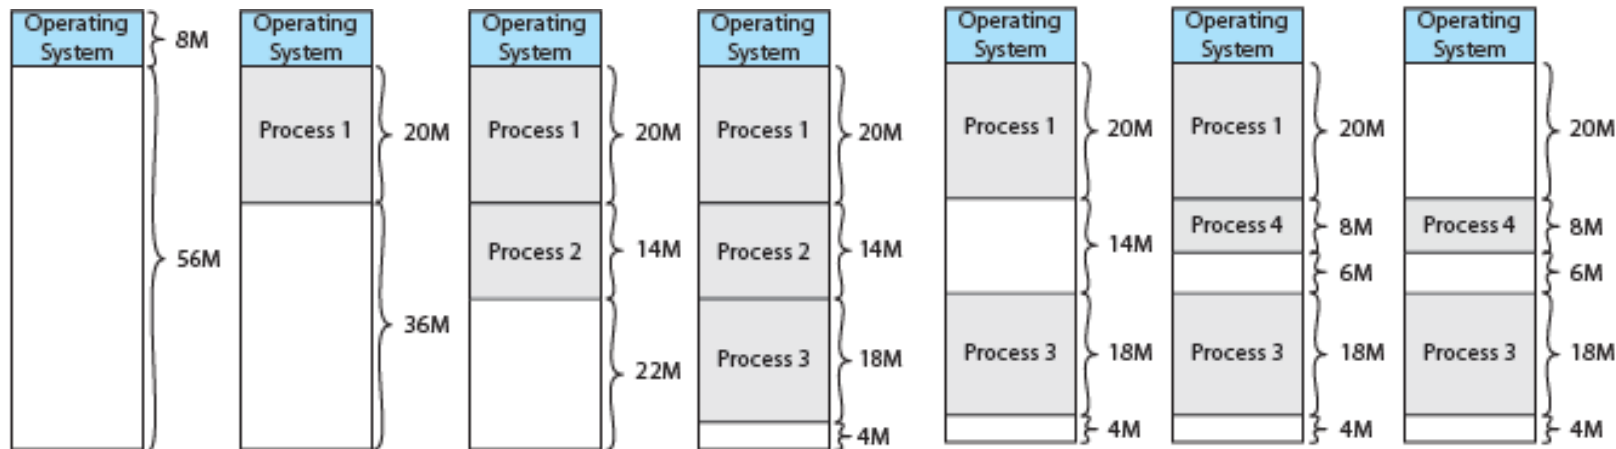
- Partitions vary in **length** & **number**
- Processes are given the **exact memory** they require
- Example (RAM 64M)
 - P1 starts (20M), P2 starts (14M), P3 starts (18M)
 - P2 ends, **P4 starts (8M)**, P1 ends, P2 restarts (14M)



Partitioning

● Dynamic Partitioning

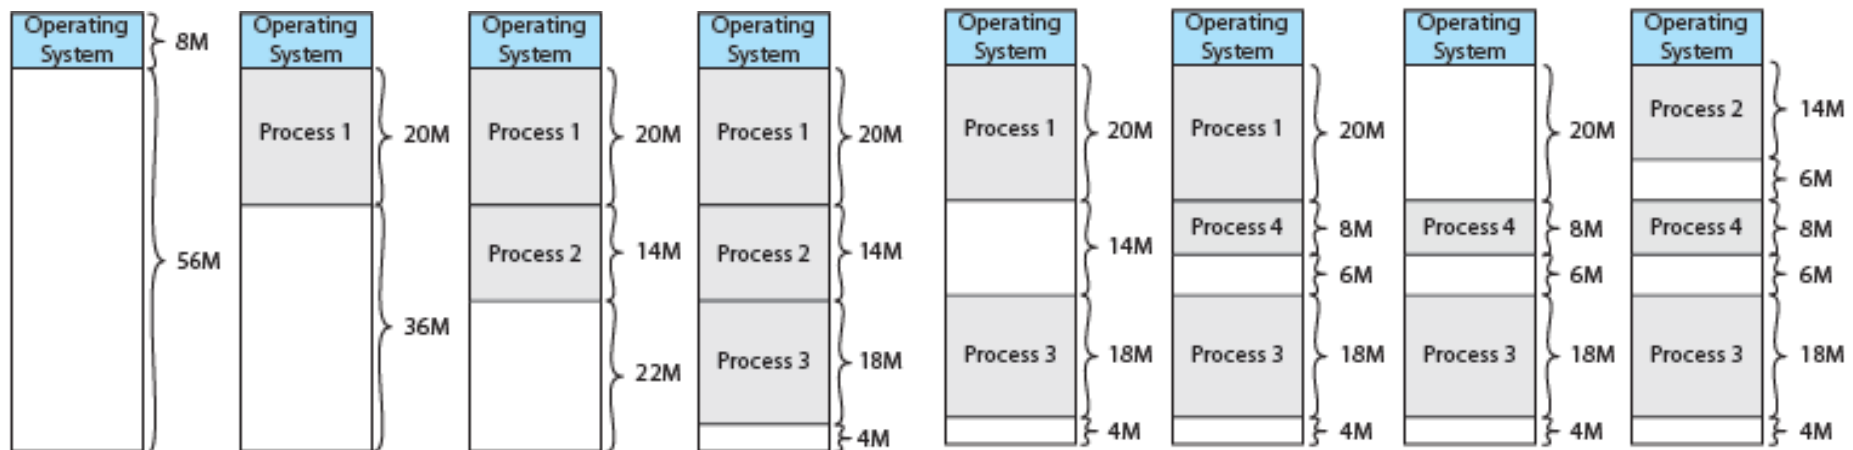
- Partitions vary in **length** & **number**
- Processes are given the **exact memory** they require
- Example (RAM 64M)
 - P1 starts (20M), P2 starts (14M), P3 starts (18M)
 - P2 ends, P4 starts (8M), **P1 ends**, P2 restarts (14M)



Partitioning

Dynamic Partitioning

- Partitions vary in **length** & **number**
- Processes are given the **exact memory** they require
- Example (RAM 64M)
 - P1 starts (20M), P2 starts (14M), P3 starts (18M)
 - P2 ends, P4 starts (8M), P1 ends, **P2 restarts (14M)**



Partitioning

Dynamic Partitioning

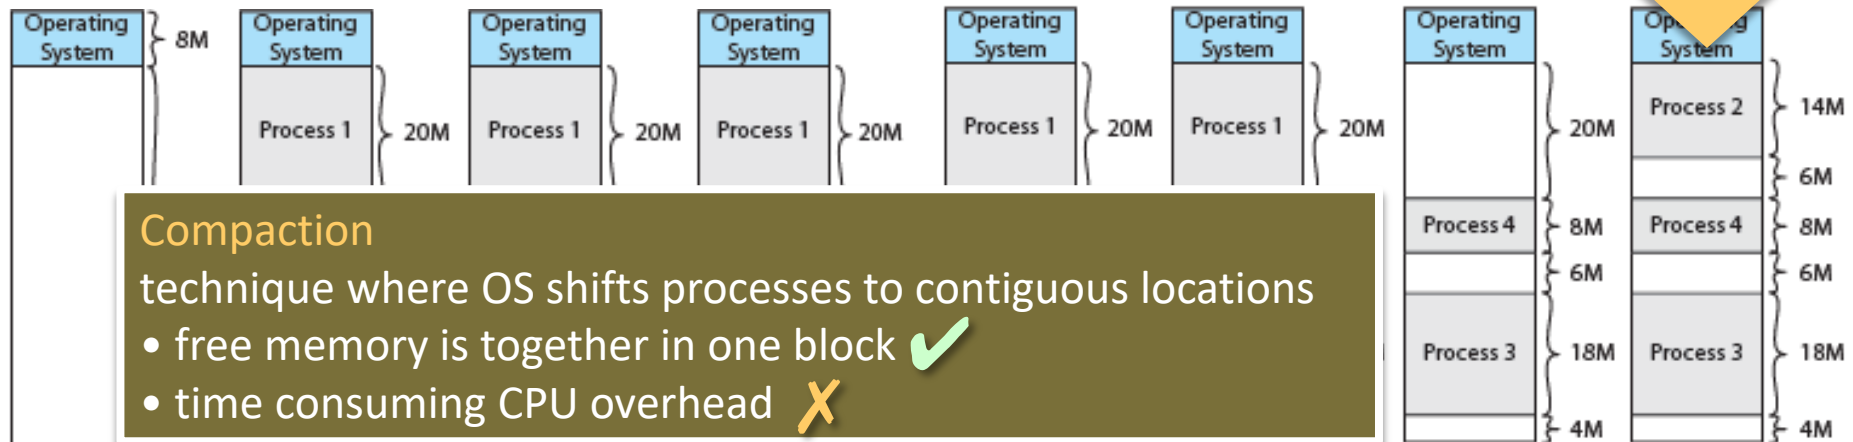
- Partitions vary in **length** & **number**
- Processes are given the **exact memory** they require

- Example (RAM 64M)

- P1 starts (20M), P2 starts (20M)
- P2 ends, P4 starts (8M)

External Fragmentation

- memory becomes more and more fragmented
- memory utilization declines



Compaction

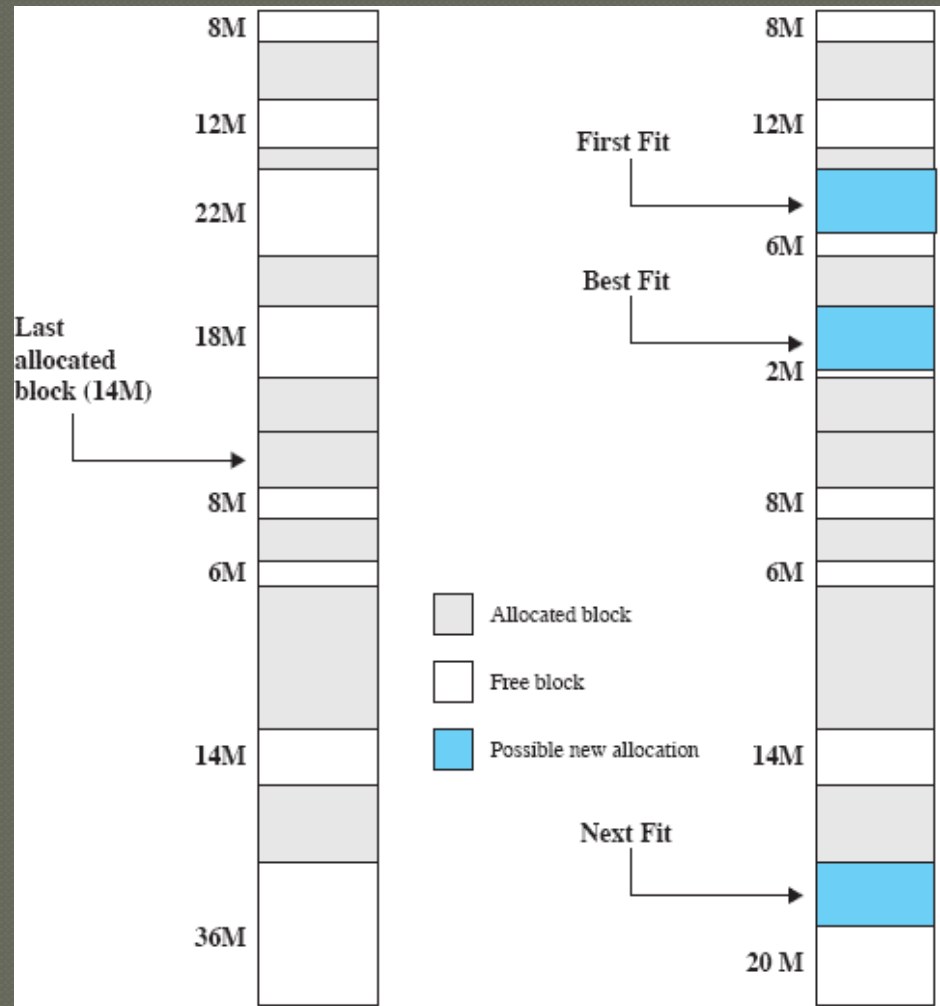
technique where OS shifts processes to contiguous locations

- free memory is together in one block ✓
- time consuming CPU overhead ✗

Partitioning

Dynamic Partitioning


- Placement algorithms
- **Best-fit**
 - chooses block **closest in size** to fit the request.
- **First-fit**
 - scanning from **top**.
 - chooses first block **large enough** to fit request
- **Next-fit**
 - scanning from place of **last allocation**.
 - chooses next block **large enough** to fit request



Partitioning

● Dynamic Partitioning

- Entire process has to be loaded in contiguous memory block
- What if you have enough memory but its fragmented?
 - Compaction



Memory Management

- ◉ Intro

- ◉ Requirements

- Relocation, Protection, Sharing, Logical & Physical organization

- ◉ Partitioning

- Fixed & Dynamic partitioning

- ◉ Paging

- Frames & pages, Addressing

Paging

- Partition **memory** into **frames**...
- Partition **processes** into **pages**...
 - ...which are **equal fixed-size** chunks relatively **small**
 - e.g., **A(4) runs**, B(3) runs, C(4) runs, B ends, D(5) runs

Frame number	Main memory		Main memory
0		0	A.0
1		1	A.1
2		2	A.2
3		3	A.3
4		4	
5		5	
6		6	
7		7	
8		8	
9		9	
10		10	
11		11	
12		12	
13		13	
14		14	

Paging

- Partition **memory** into **frames**...
- Partition **processes** into **pages**...
 - ...which are **equal fixed-size** chunks relatively **small**
 - e.g., A(4) runs, **B(3) runs**, C(4) runs, B ends, D(5) runs

Frame number	Main memory	Main memory	Main memory
0		A.0	A.0
1		A.1	A.1
2		A.2	A.2
3		A.3	A.3
4			B.0
5			B.1
6			B.2
7			
8			
9			
10			
11			
12			
13			
14			

Paging

- Partition **memory** into **frames**...
- Partition **processes** into **pages**...
 - ...which are **equal fixed-size** chunks relatively **small**
 - e.g., A(4) runs, B(3) runs, **C(4) runs**, B ends, D(5) runs

Frame number	Main memory	Main memory	Main memory	Main memory
0		A.0	A.0	A.0
1		A.1	A.1	A.1
2		A.2	A.2	A.2
3		A.3	A.3	A.3
4			B.0	B.0
5			B.1	B.1
6			B.2	B.2
7				C.0
8				C.1
9				C.2
10				C.3
11				
12				
13				
14				

Paging

- Partition **memory** into **frames**...
- Partition **processes** into **pages**...
 - ...which are **equal fixed-size** chunks relatively **small**
 - e.g., A(4) runs, B(3) runs, C(4) runs, **B ends**, D(5) runs

Frame number	Main memory	Main memory	Main memory	Main memory	Main memory
0		A.0	A.0	A.0	A.0
1		A.1	A.1	A.1	A.1
2		A.2	A.2	A.2	A.2
3		A.3	A.3	A.3	A.3
4			B.0	B.0	
5			B.1	B.1	
6			B.2	B.2	
7				C.0	C.0
8				C.1	C.1
9				C.2	C.2
10				C.3	C.3
11					
12					
13					
14					

Paging

- Partition **memory** into **frames**...
- Partition **processes** into **pages**...
 - ...which are **equal fixed-size** chunks relatively **small**
 - e.g., A(4) runs, B(3) runs, C(4) runs, B ends, **D(5) runs**

Frame number	Main memory	Main memory	Main memory	Main memory	Main memory	Main memory
0		A.0	A.0	A.0	A.0	A.0
1		A.1	A.1	A.1	A.1	A.1
2		A.2	A.2	A.2	A.2	A.2
3		A.3	A.3	A.3	A.3	A.3
4			B.0	B.0		D.0
5			B.1	B.1		D.1
6			B.2	B.2		D.2
7				C.0	C.0	C.0
8				C.1	C.1	C.1
9				C.2	C.2	C.2
10				C.3	C.3	C.3
11						D.3
12						D.4
13						
14						

Partition m

Partition p

...which are

e.g., A(4) r

Page Table

- Table where OS keeps frame location of each process page
- Used by processor to produce a physical address

How?

0	0
1	1
2	2
3	3

Process A
page table

0	—
1	—
2	—

Process B
page table

0	7
1	8
2	9
3	10

Process C
page table

0	4
1	5
2	6
3	11
4	12

Process D
page table

13
14

Free frame
list




Frame number	Main memory	Main memory	Main memory	Main memory	Main memory	Main memory
0		A.0	A.0	A.0	A.0	A.0
1		A.1	A.1	A.1	A.1	A.1
2		A.2	A.2	A.2	A.2	A.2
3		A.3	A.3	A.3	A.3	A.3
4			B.0	B.0		D.0
5			B.1	B.1		D.1
6			B.2	B.2		D.2
7				C.0	C.0	C.0
8				C.1	C.1	C.1
9				C.2	C.2	C.2
10				C.3	C.3	C.3
11						D.3
12						D.4
13						
14						

Paging – (fixed page size)

● Addressing: break up into frames


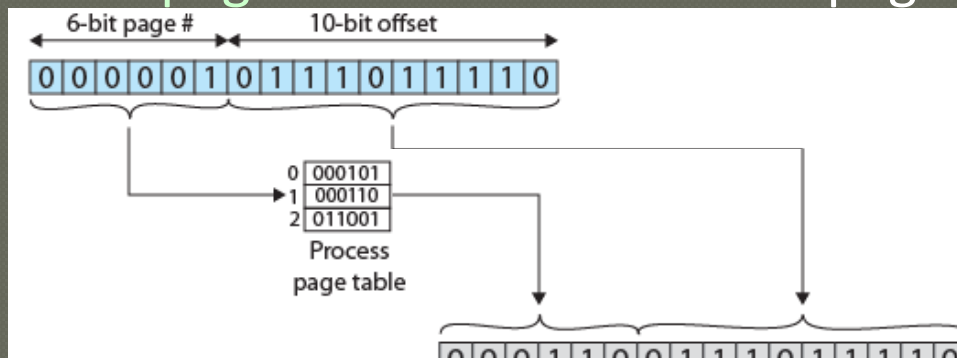
- Divide address into page bits and address bits
- Length of each page? 2^{**} address bits
- Number pages? 2^{**} page bits
 - Ex. if 16-bit addressing & 1K page size
10 bits page length $\Rightarrow 2^{10} = 1024$
6 bits page number $\Rightarrow 2^6 = 64$ pages



Paging – (fixed page size)

Addressing: from relative to physical

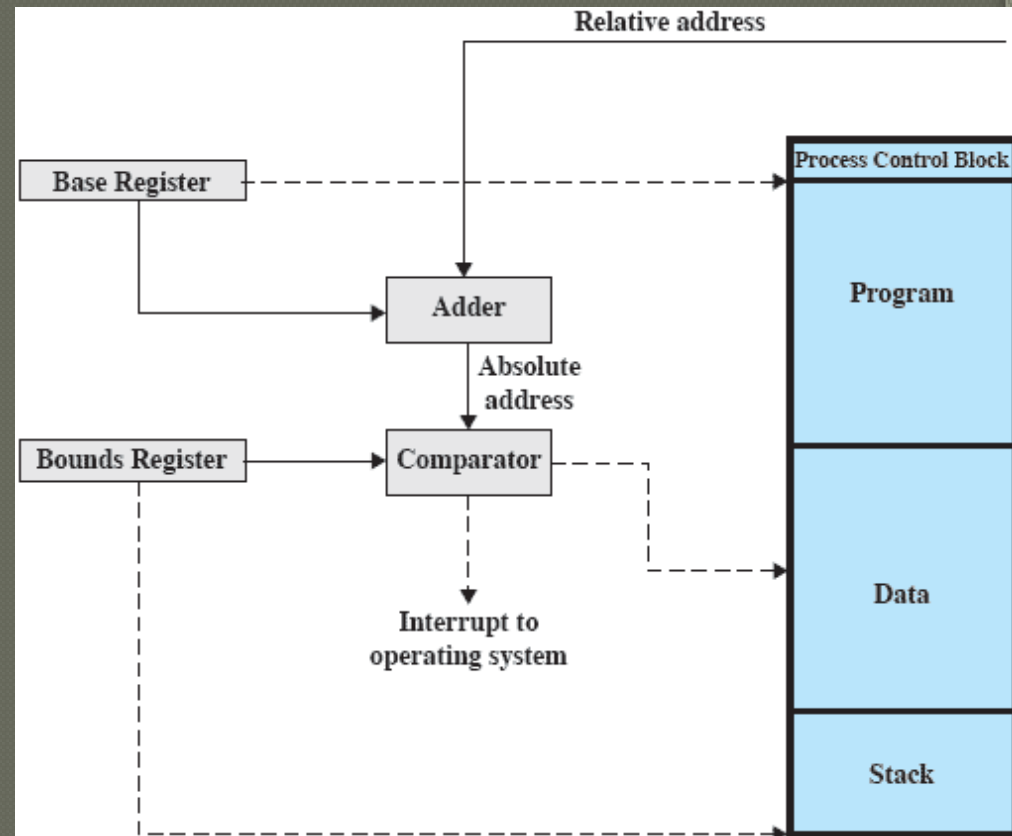
- Divide address into page bits and address bits
- Length of each page? 2^{**} address bits
- Number pages? 2^{**} page bits
 - Ex. if 16-bit addressing & 1K page size
10 bits page length $\Rightarrow 2^{**}10 = 1024$
6 bits page number $\Rightarrow 2^{**}6 = 64$ pages



Partitioning

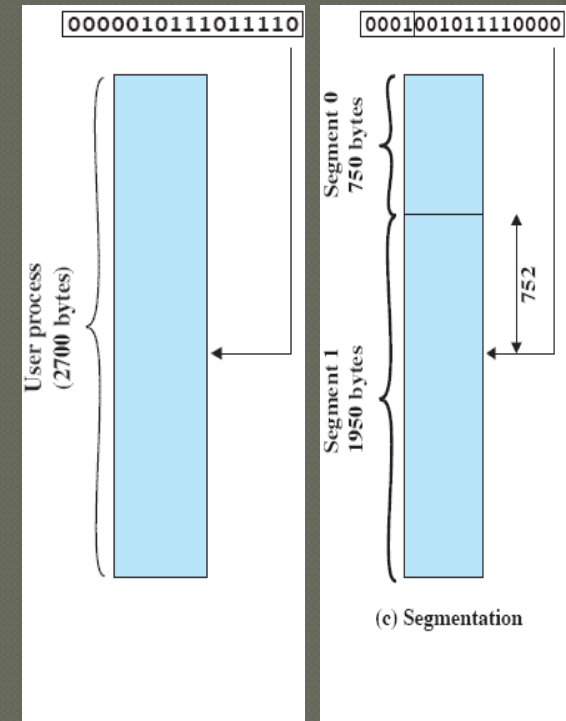
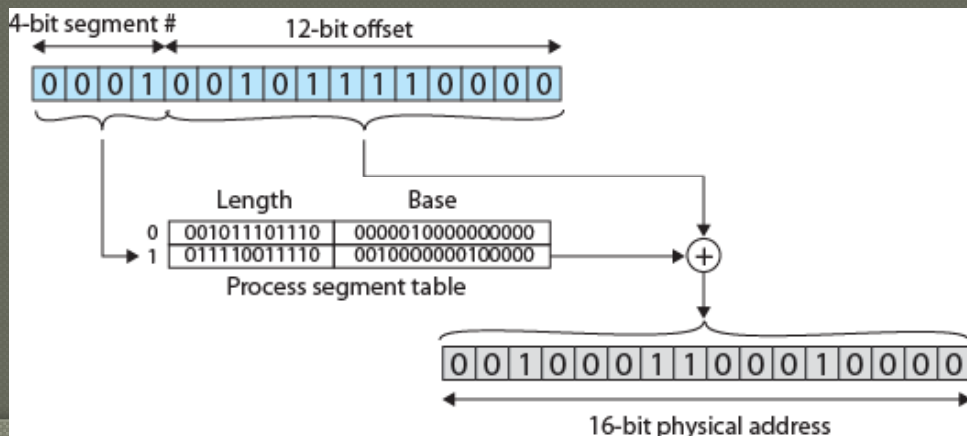
Addresses

- Logical
 - reference to a memory location **independent** of the current assignment of data to memory
- Relative
 - (special case of logical) reference to a memory location **relative** to a known point
- Physical (aka absolute)
 - reference to the **actual** location in main memory



Partitioning

- Partition processes into **segments**...
 - chunks of **varying** but **limited length**
- Same addressing layout applies:
 - segment number + offset
 - e.g., if **16-bit** addressing & **12 bits** for reference anywhere then there can be **32 segments** of up to **4K** max in length



Memory Management

- Intro

- Requirements

- Relocation, Protection, Sharing, Logical & Physical organization

- Partitioning

- Fixed & Dynamic Partitioning, Buddy system

- Paging

- Frames & pages, Addressing

Done!